



查询处理与查询优化

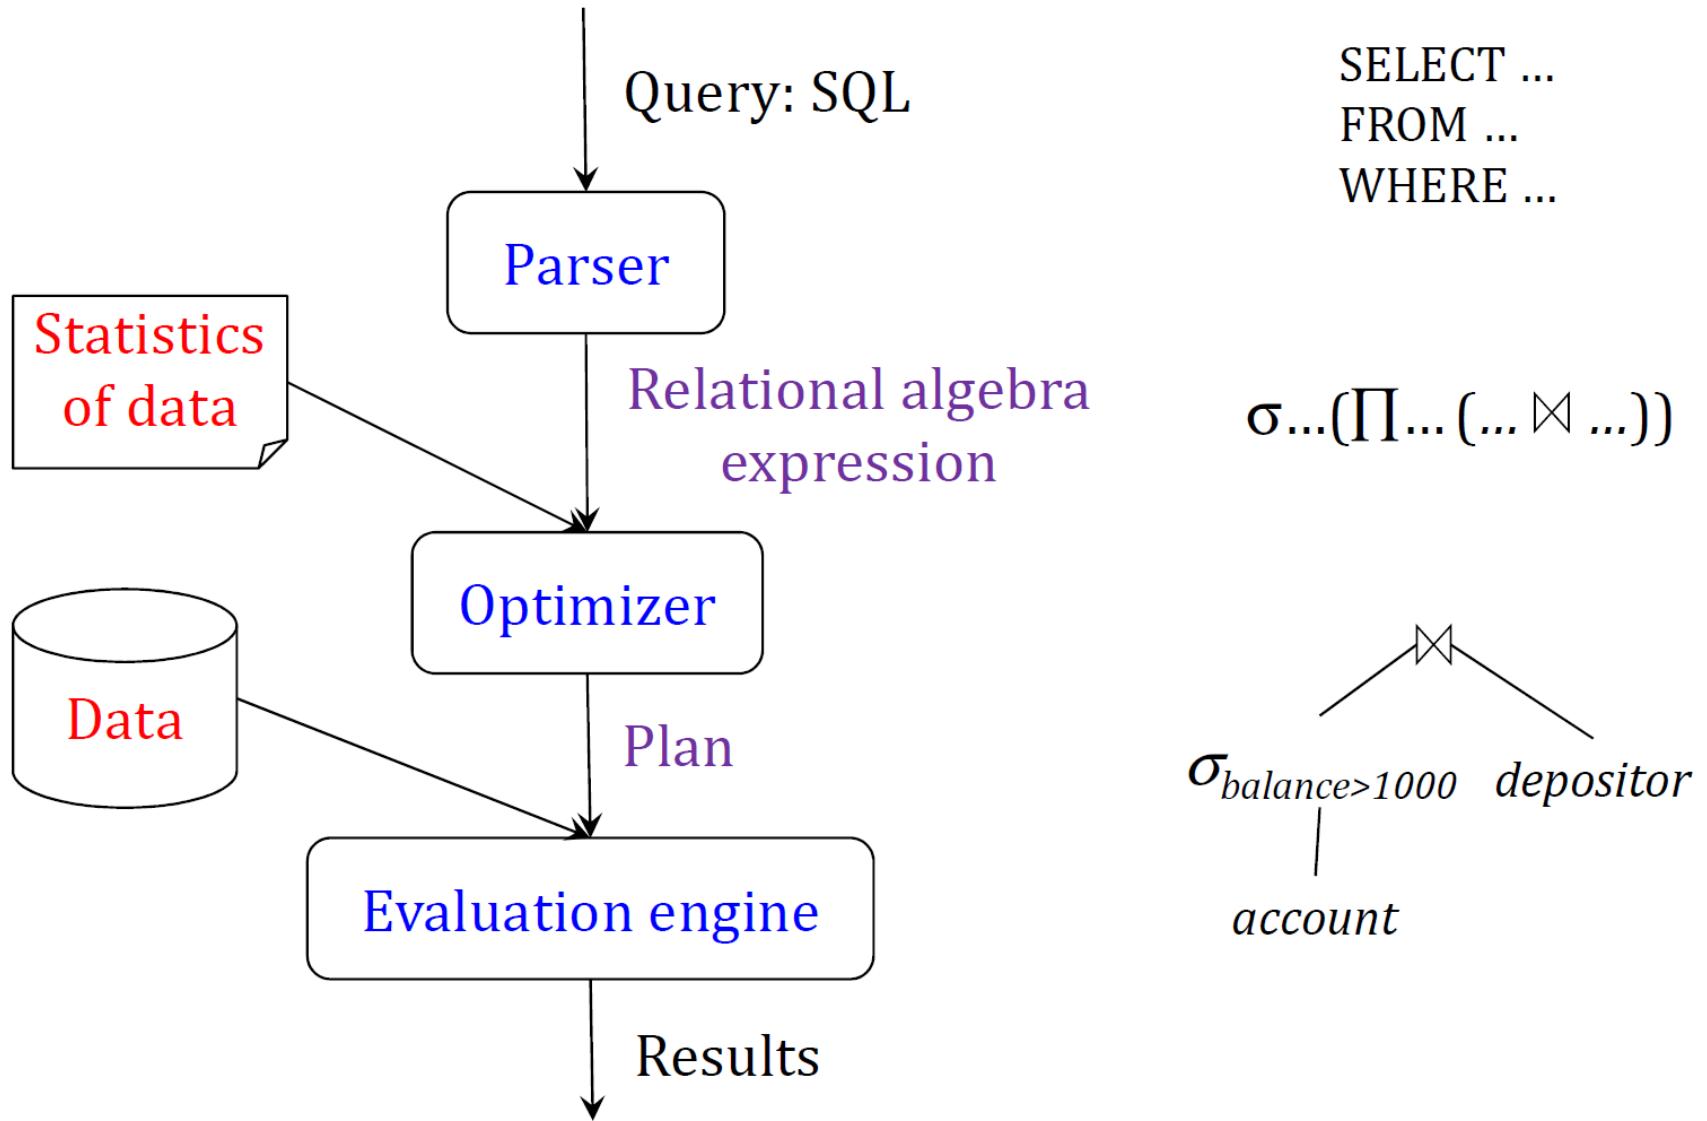
南方科技大学

唐 博

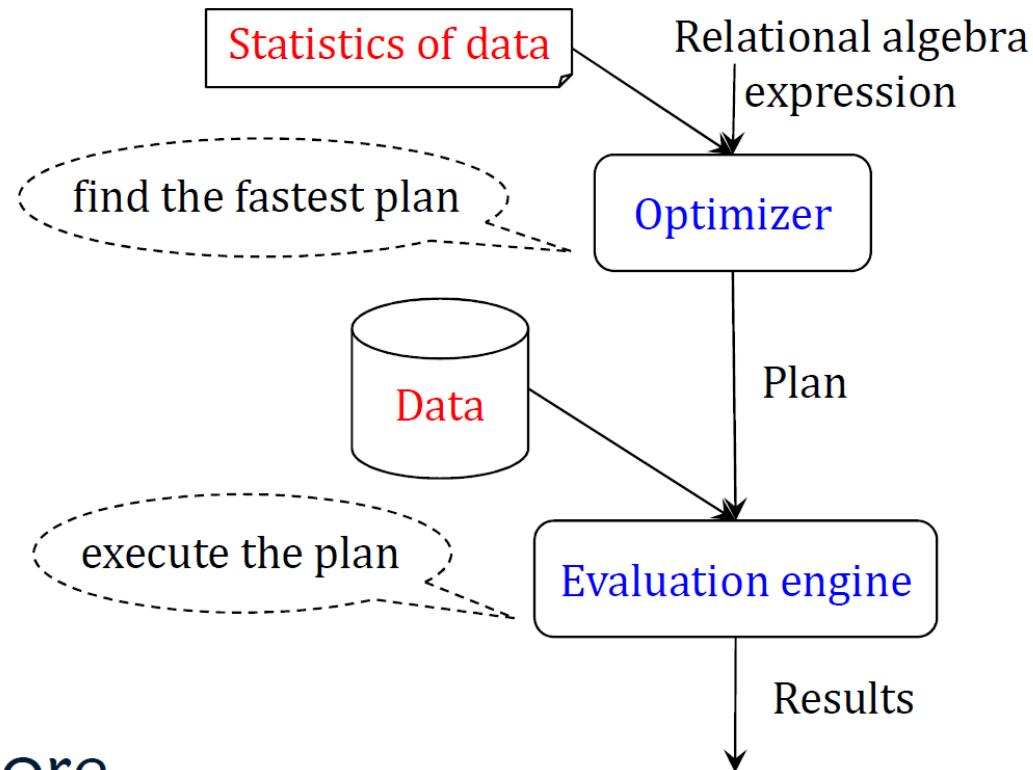
tangb3@sustech.edu.cn



How to process a SQL query

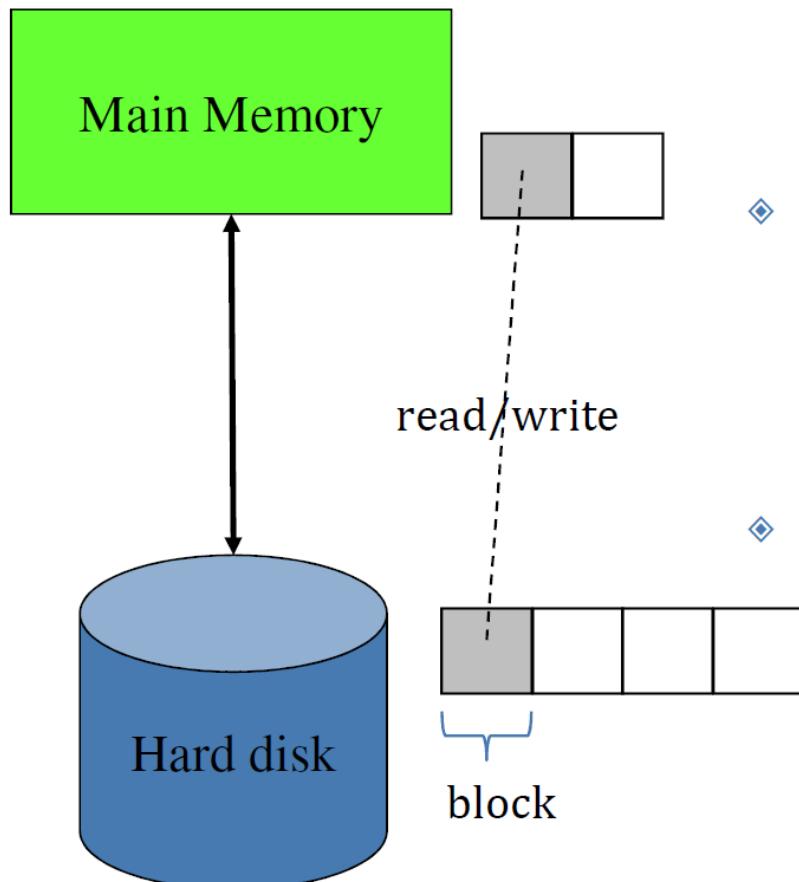


How to process a SQL query



Issues to explore

1. How to store data in the hard disk?
2. Evaluation engine: How to execute a plan?
3. Optimizer: How to find the fastest plan?



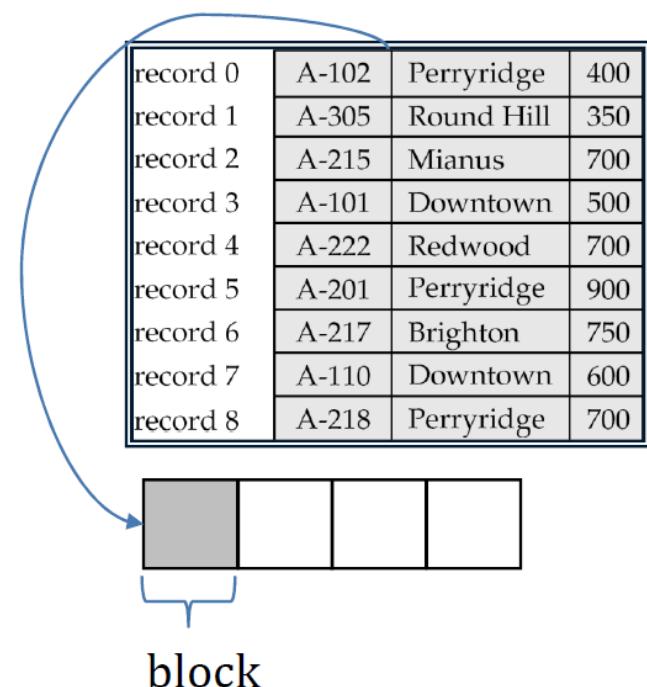
- ❖ RDBMS stores data in hard disk because
 - ❖ Large capacity (e.g., 10 TB)
 - ❖ Non-volatile (Keep data without power)
- ❖ **Disk block** – unit of data transfer between disk and main memory
 - ❖ Typical block size: 4 kB
 - ❖ Much larger than an attribute's size
- ❖ Access time = $b * t_T + s * t_S$
 - b – the number of disk block transfers
 - s – the number of seeks
 - t_T – time to transfer one disk block
 - t_S – time to seek one disk block
- ❖ Typical values: seek time = 4-10 ms, data transfer rate = 25-100 MB/s
 - much slower than CPU time!

File Organization

- ❖ How to store a relation (as a file) in the hard disk?
 - ❖ File = a sequence of *records*
 - ❖ Record = a sequence of fields (attribute values)
 - ❖ Assume fixed-length record

- ❖ Notations for a relation r
 - ❖ n_r : # of records in r
 - ❖ f_r : # of records per disk block
 - ❖ b_r : # of blocks

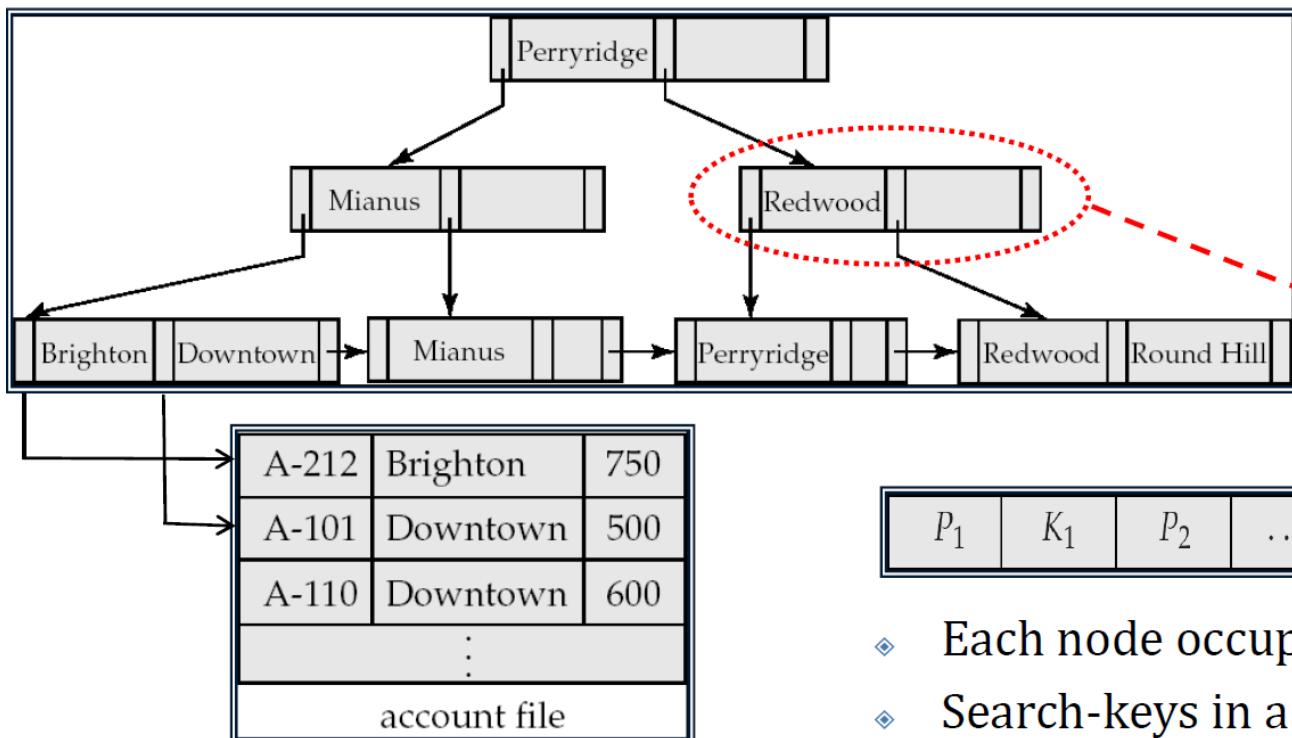
$$b_r = n_r / f_r$$



- ❖ How to search records efficiently?

B⁺-tree Structure

- ❖ Except the root node, each node must be at least **half full**
- ❖ Non-leaf nodes vs. leaf nodes
 - ❖ *Sibling pointer* in leaf node is used to process range queries



B⁺-tree for *account* file

B⁺-tree node



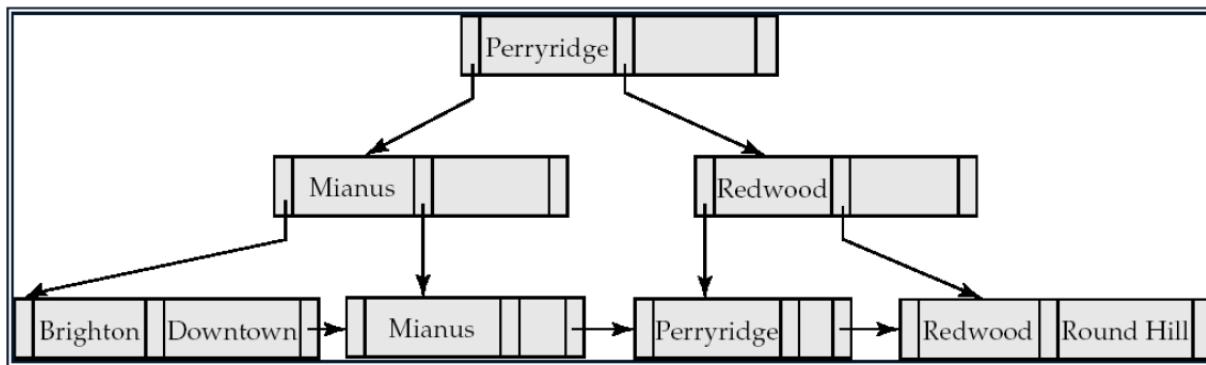
- ❖ Each node occupies a disk **block**
- ❖ Search-keys in a node are **ordered** :

$$K_1 < K_2 < K_3 < \dots < K_{n-1}$$
- ❖ Each pointer is a disk block address

- ❖ How to find records with a search-key value of k ?
 1. Start with the *root node*
 2. Recursively follow a pointer so that it is smallest search-key $> k$
 3. When at leaf node, find the pointer to the desired record (if exists)

Observations

- ❖ Lookup cost = tree height = $\lceil \log_{f/2}(K) \rceil$
 - ❖ K : # of search-key values in the file
 - ❖ f : maximum # of pointers in a node
- ❖ Example: given $K = 1$ million,
 block size = 4096 bytes,
 key size = 40 bytes
 pointer size = 4 bytes
 - ❖ How to calculate f ? $f = 94$
 - ❖ Lookup cost
 $= \lceil \log_{47}(1,000,000) \rceil = 4$ blocks

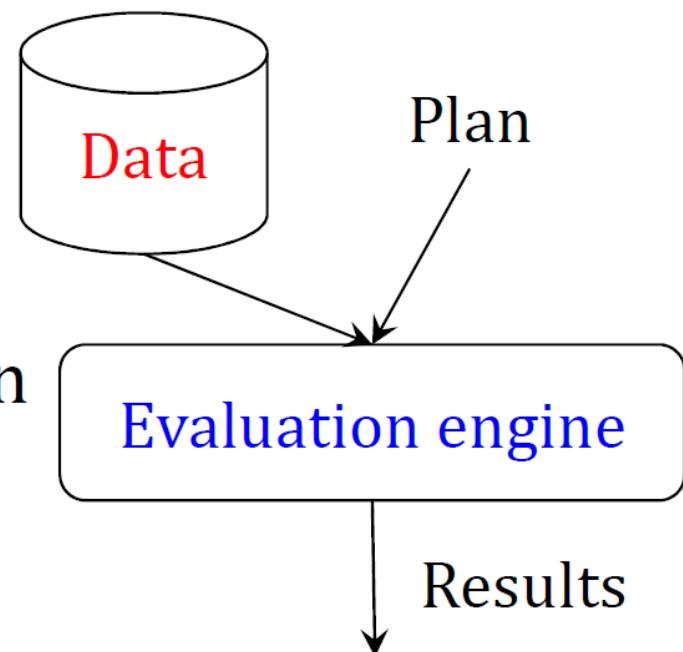


Evaluation Engine

- ❖ How to measure the cost of a plan?

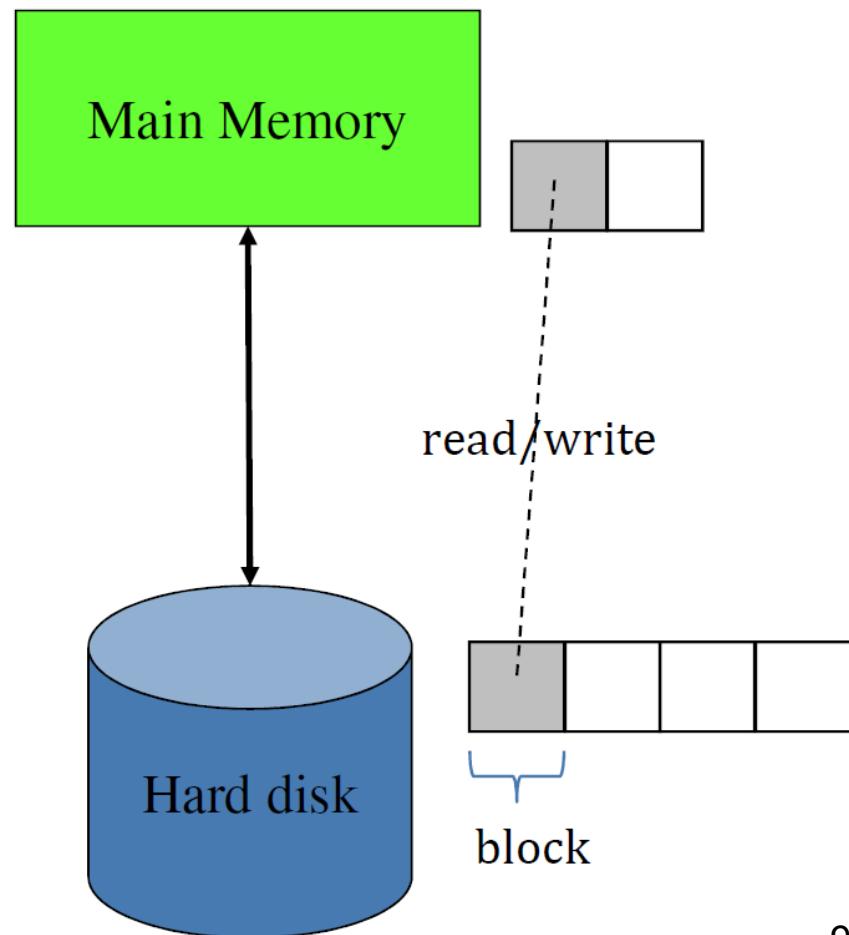
- ❖ Some methods for executing a selection operator

- ❖ Some methods for executing a join operator



Execution Cost

- ❖ Cost = the number of disk block transfers
- ❖ Assumptions in RDBMS
 - ❖ Ignore CPU costs
 - ❖ Ignore the cost of writing the final output to disk
- ❖ Extra assumption in this course
 - ❖ Ignore the disk seek time, because the number of disk block transfers \gg the number of disk seeks





Selection Operation



- ❖ Example: $\sigma_{\text{balance} < 2500}(\text{account})$
- ❖ Several different algorithms to implement selections
 - ❖ Usually choose the cheapest available one

Algorithm / physical operator	Cost (# disk blocks)
Linear search	b_r
Primary index, equality on key	$h_r + 1$
Primary index, equality on non-key	$h_r + b_{\text{results}}$
Secondary index, equality	$h_r + n_{\text{results}}$

b_r : size of r in blocks

h_r : height of B⁺-tree on r

n_{results} : # of results

b_{results} : result size in blocks

For simplicity, we measure the cost as the number of disk block transfers.

- ❖ *Linear search:*

Scan each file block and test all tuples

- ❖ Applicable to any type of condition
- ❖ Cost = b_r blocks (number of blocks occupied by relation r)

- ❖ *Primary index on candidate key, equality:*

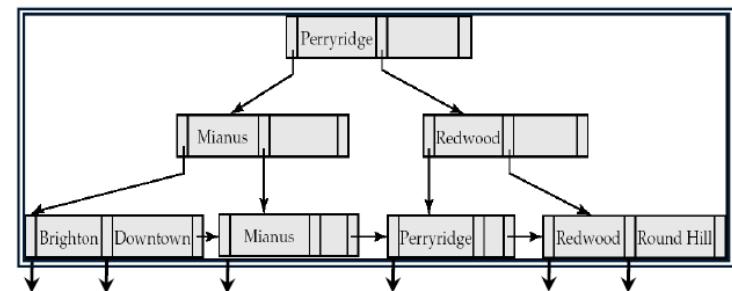
Retrieve a single tuple that satisfies the equality condition

- ❖ Cost = $h_r + 1$ blocks
 - ❖ B⁺-tree height $h_r = \lceil \log_{f/2} n_r \rceil$

- ❖ *Primary index on nonkey, equality:*

Retrieve multiple (consecutive) tuples that satisfies

- ❖ Cost = $h_r + b_{results}$ blocks
 - ❖ Let $b_{results}$ = number of blocks containing matching tuples





Join Operation



- ❖ Example: $customer \bowtie depositor$
- ❖ Several different algorithms to implement joins
 - ❖ Usually choose the cheapest available one

Algorithm / physical operator	Cost (# disk blocks)
Nested-loop join	$n_r * b_s + b_r$ [worst case]
Block nested-loop join	$\lceil b_r / (M-2) \rceil * b_s + b_r$
Indexed nested-loop join	$n_r * (h_s + 1) + b_r$
Merge-join	$b_r + b_s$ $+ b_r (2 \lceil \log_{M-1} (b_r/M) \rceil + 1)$ $+ b_s (2 \lceil \log_{M-1} (b_s/M) \rceil + 1)$
Hash-join	$3(b_r + b_s) + 4 n_h$ if no recursive partitioning required

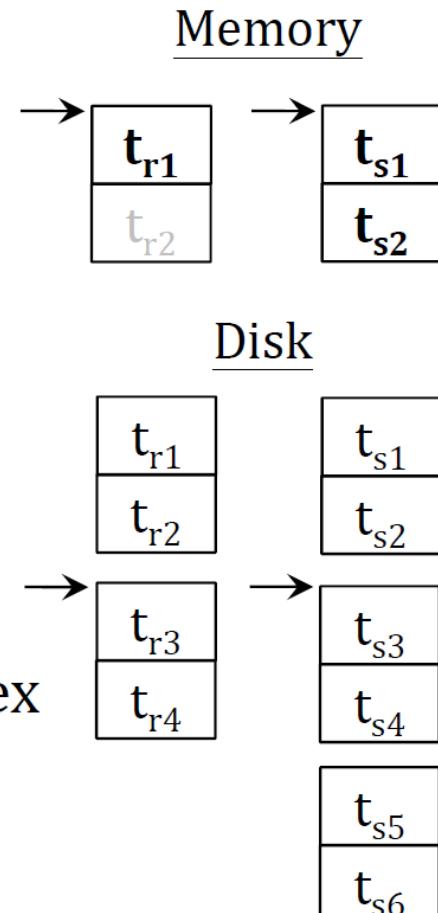
For simplicity, we measure the cost as the number of disk block transfers.

Nested-Loop Join

```

for each tuple  $t_r$  in  $r$  do
    Outer relation
    for each tuple  $t_s$  in  $s$  do
        Inner relation
        if pair  $(t_r, t_s)$  satisfies the join condition  $\theta$ 
        then add  $t_r \cdot t_s$  to the result
    
```

$$r \bowtie_{\theta} s$$



- ◊ Applicable to any join condition, requires no index
- ◊ But expensive!
- ◊ Improvement: *Index Nested-Loop Join*
 - ◊ Use index on inner relation if available



Nested-Loop Join



- ❖ Cost of *nested loop join*: $n_r * b_s + b_r$ blocks
 - ❖ Assume the worst case: only one memory buffer block for each relation
- ❖ Cost of *indexed nested loop join*: $n_r * (h_s + 1) + b_r$ blocks
- ❖ **Exercise:** $customer \bowtie depositor$
 - ❖ Number of tuples: $n_{customer} = 10000$ $n_{depositor} = 5000$
 - ❖ Number of blocks: $b_{customer} = 400$ $b_{depositor} = 100$
 - ❖ Suppose that *customer* has a primary B⁺-tree index on the join attribute *customer-name*, which contains 20 entries per index node.
 - ❖ Cost of *nested loop join*

If outer relation = *depositor* : _____ * _____ + _____ = _____ blocks
If outer relation = *customer* : _____ * _____ + _____ = _____ blocks
 - ❖ Cost of *indexed nested loop join*

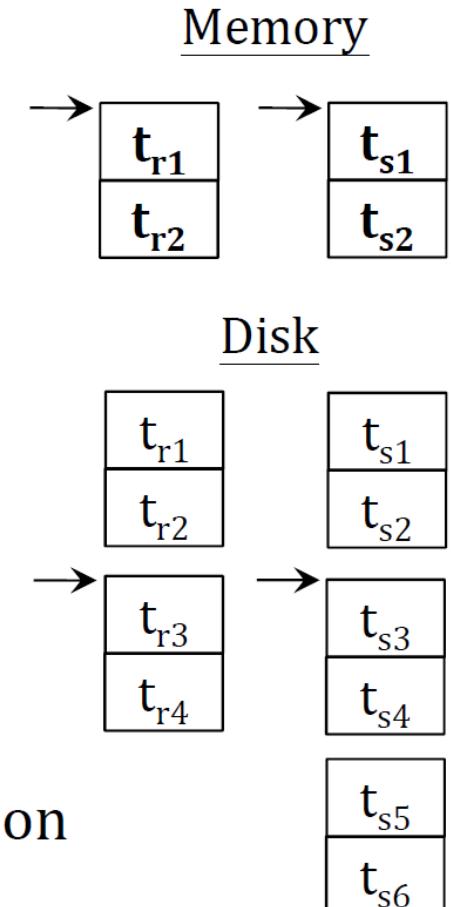
E.g., the tree height (for *customer*) is 4, so the cost is:
 $5000 * (4+1) + 100 = 25,100$ blocks

$$\begin{aligned}& \lceil \log_{f/2}(n_{customer}) \rceil \\&= \lceil \log_{20/2}(10000) \rceil = 4\end{aligned}$$

Block Nested-Loop Join

- ❖ Variant of nested-loop join
- ❖ Pair every block of the inner relation with every block of the outer relation

```
for each block  $B_r$  of  $r$  do
    for each block  $B_s$  of  $s$  do
        for each tuple  $t_r$  in  $B_r$  do
            for each tuple  $t_s$  in  $B_s$  do
                if pair  $(t_r, t_s)$  satisfies the join condition
                then add  $t_r \cdot t_s$  to the result
```





Block Nested-Loop Join

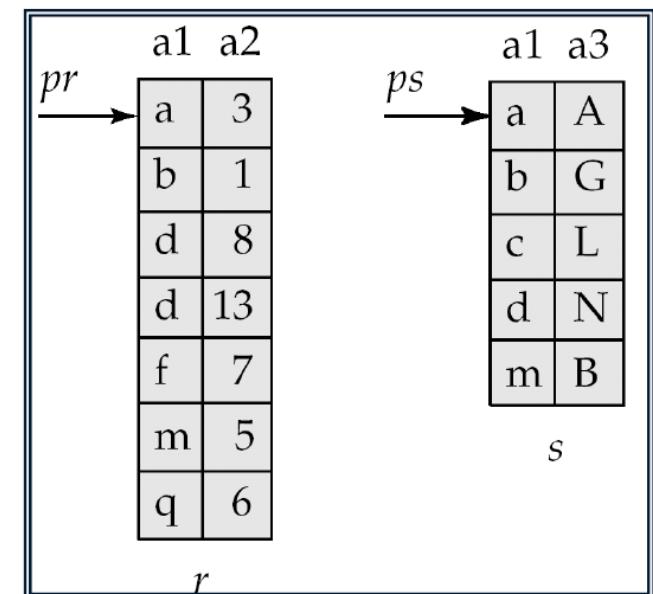


- ❖ Worst case cost: $b_r * b_s + b_r$ blocks
 - ❖ Assume the worst case: only one memory buffer block for each relation
- ❖ If the memory has M blocks,
$$\text{Cost} = \lceil b_r / (M-2) \rceil * b_s + b_r \text{ blocks}$$
 - ❖ Use the memory blocks as follows:
 - $M - 2$ memory blocks to buffer the outer relation
 - 1 block to buffer the inner relation
 - 1 block to buffer the output

Merge-Join

1. Sort both relations on their join attribute (*to discuss soon*)
2. Merge the sorted relations to join them
 - a. This step is similar to the merge stage of the sort-merge algorithm.
 - b. The difference is to handle duplicate values in join attribute
 - every pair with same value on join attribute must be matched

- ◊ Applicable to equi-joins and natural joins
- ◊ Each block is only read once
 - ◊ assuming all tuples for any given value of the join attributes fit in memory
- ◊ Cost of merge join =
 $b_r + b_s$ blocks
 - + the sorting cost (if relations are unsorted)



External Sort Merge-Join

- ❖ Use it when the relation is larger than the main memory, i.e., $b_r > M$
- ❖ **External sort-merge algorithm**
 - ❖ 1. Create sorted runs
 - ❖ Read consecutive M blocks into memory, sort it, then write to a run
 - ❖ 2. Merging until only 1 run remains
- ❖ Cost: $b_r (2 \lceil \log_{M-1}(b_r/M) \rceil + 1)$
 - ❖ Number of merge passes required: $\lceil \log_{M-1}(b_r/M) \rceil$
 - ❖ Block transfers for initial run and in each pass is $2b_r$

Example: $M = 3$
memory blocks

a 19		
d 31		
g 24		
b 14		
c 33		
b 14		
c 33		
e 16		
r 16		
d 21		
m 3		
r 16		
a 14		
d 7		
p 2		
initial relation		runs
create runs		

External Sort Merge-Join

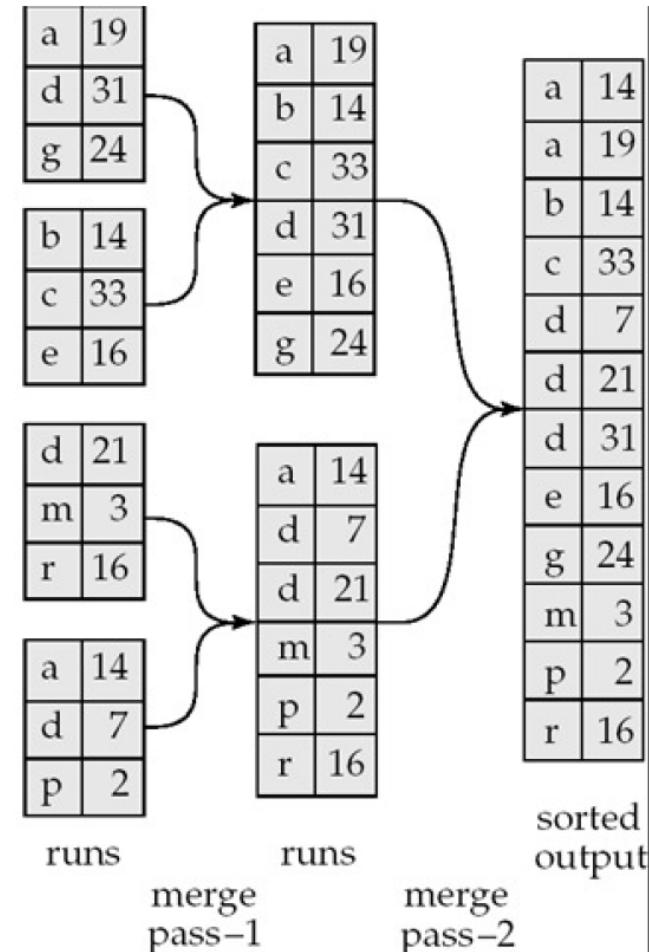
- How to merge “sorted runs”?

2. Merging (for every $M-1$ runs)

- Use $M-1$ memory blocks as input buffers, and 1 memory block as output buffer
- Move the smallest tuple from its input buffer to the output buffer
- An input block empty → fill it with the next disk block from its input run
- An output block full → flush the block to its output run

- Repeat until only 1 run remains

Example: $M = 3$ memory blocks



Hash Join Algorithm

- * Applicable to equi-joins and natural joins

1. Partition the relation s using a **hash function h** (into n_h partitions)

* note that n_h must be smaller than M

2. Partition r similarly

3. For each partition value i

- (a) Create another hash function h'

- (b) Load s_i into memory.

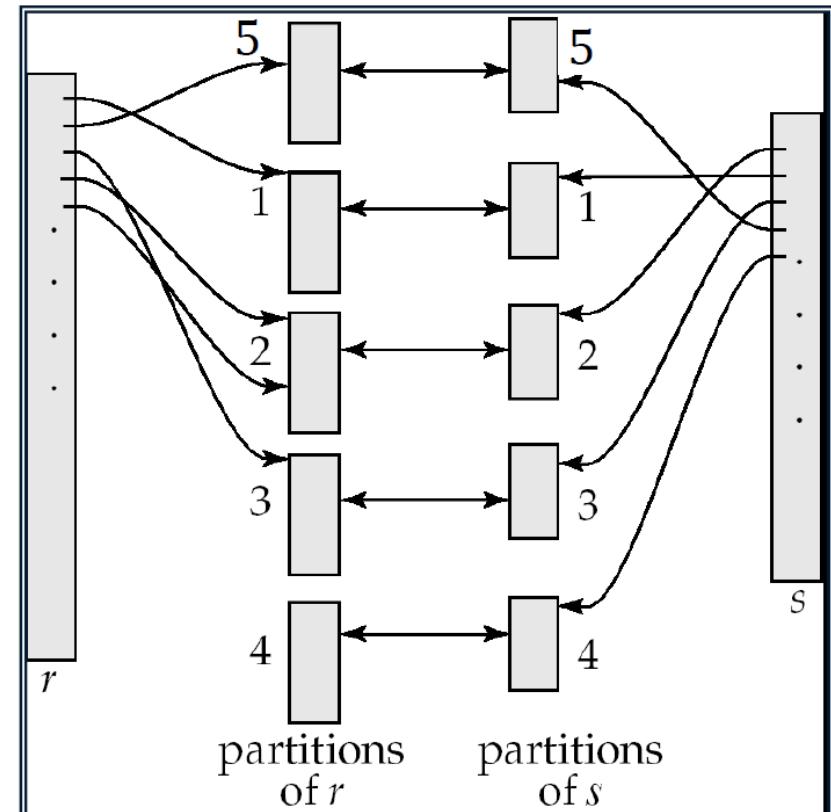
Build an in-memory hash index on it by the join attribute (by using h').

- (c) Read the tuples in r_i .

For each tuple t_r , find each matching tuple t_s in s_i from in-memory hash index.

Output the result.

Example: $n_h = 5$ partitions



r : probe input

s : build input

Hash Join Algorithm

- ❖ Suppose that number of partitions n_h is at most number of memory blocks M
- ❖ Cost of hash join:
 - ❖ Partitioning phase: $3(b_r + b_s) + 4 n_h$
 - ❖ Build and probe phase: $2(b_r + b_s)$
 - ❖ Partially filled blocks: $b_r + b_s$
 - 2 ($2 n_h$)
- small, can be ignored
- ❖ **Recursive partitioning** is required if $n_h > M$
 - ❖ Rarely required: e.g., recursive partitioning not needed for relations of 1GB with memory size of 2MB and block size of 4KB



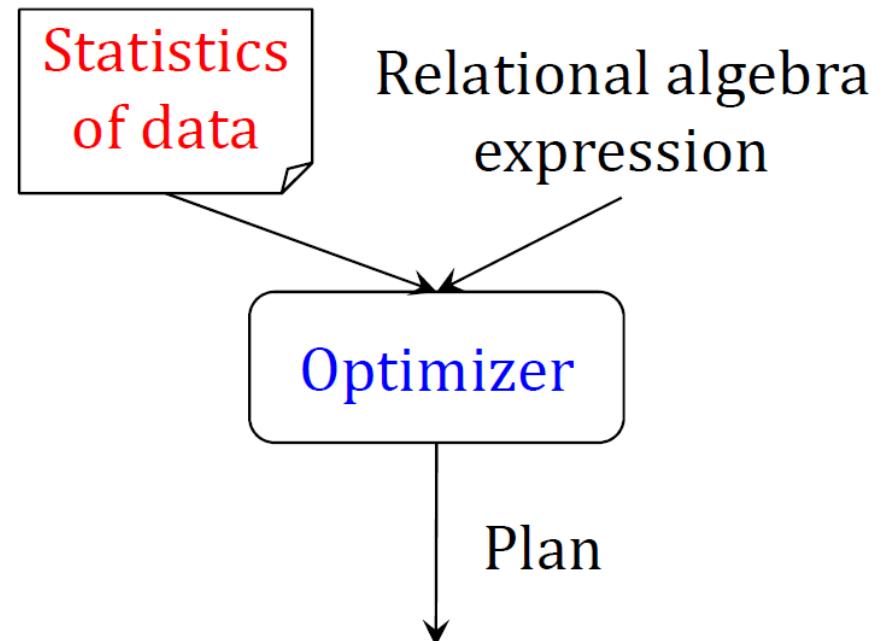
Cost of Hash Join



customer \bowtie *depositor*

- ◆ Given that
 - ◆ memory size is 20 blocks
 - ◆ $b_{depositor} = 100$ and $b_{customer} = 400$.
- ◆ Use the smaller relation (*depositor*) as build input
- ◆ How large should a partition be?
 - ◆ To make each partition of *depositor* fit in memory (20 blocks), we can partition it into $100/20 = 5$ partitions
 - ◆ Since $5 < 20$, this partitioning can be done in **one pass**
 - ◆ Similarly, we partition *customer* into 5 partitions
- ◆ Therefore total cost:
 - ◆ $3(100 + 400) = 1500$ blocks

- ❖ Why optimizer?
- ❖ Techniques
 - ❖ Equivalence rules
 - ❖ Join ordering by dynamic programming
 - ❖ Cost and size estimation

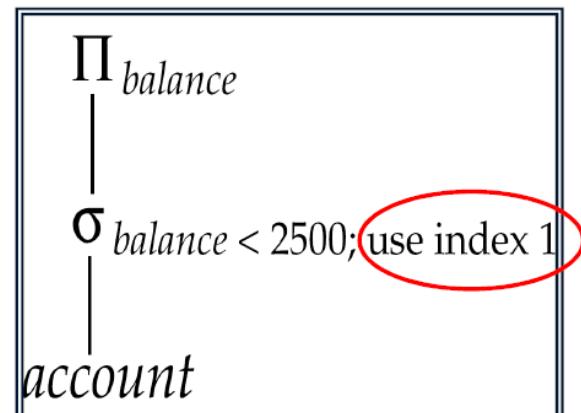




Query Optimizer



- ❖ A query (in relational algebra) has equivalent expressions
 - ❖ $\sigma_{balance < 2500}(\Pi_{balance}(account))$
 - ❖ $\Pi_{balance}(\sigma_{balance < 2500}(account))$
- ❖ An operation (e.g., selection $\sigma_{balance < 2500}$) can be evaluated by different algorithms
 - ❖ Scan the whole *balance* relation
 - ❖ Search an index on *balance*
- ❖ Different plans can have very different costs (e.g. seconds vs. days)





Why Query Optimizer?



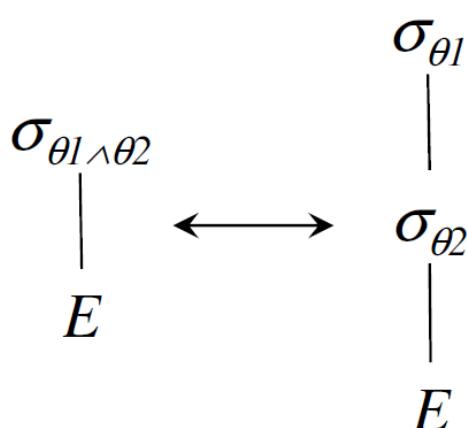
- ❖ A query optimizer executes two steps
- ❖ 1. Query rewriting
 - ❖ Relational algebra expression → logical plan
 - ❖ Rewrite a query by using materialized views
 - ❖ Obtain a cheaper plan by using **equivalence rules**
 - ❖ Perform selection early (reduces the number of tuples)
 - ❖ Perform projection early (reduces the number of attributes)
- ❖ 2. Cost-based enumeration
 - ❖ Logical plan → Physical plan
 - ❖ Enumerate join ordering by using **dynamic programming**
 - ❖ Estimate the cost of each candidate plan, then pick the cheapest plan
 - ❖ Need to **estimate the cost and output size** of operators

Equivalence Rules

- ◆ **Equivalence rule:** two relational algebra expressions output the same set of tuples on every database instance

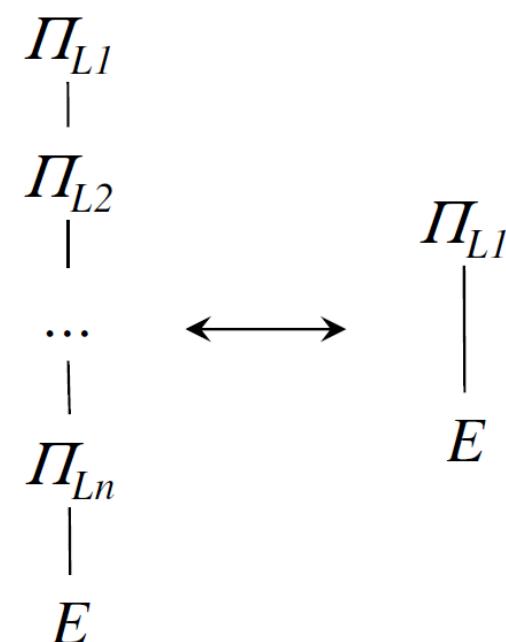
Conjunctive selection operations
can be decomposed

$$\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$



Only need the last project operation

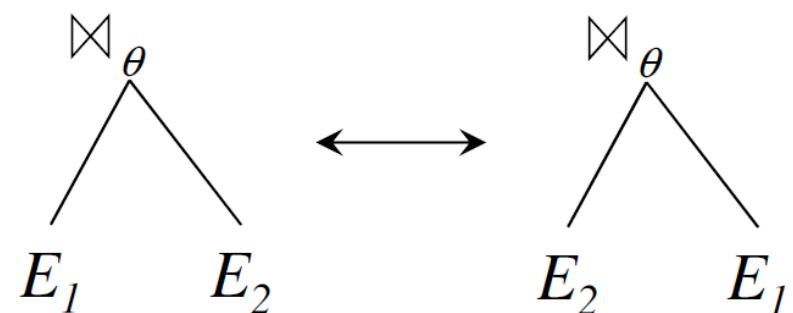
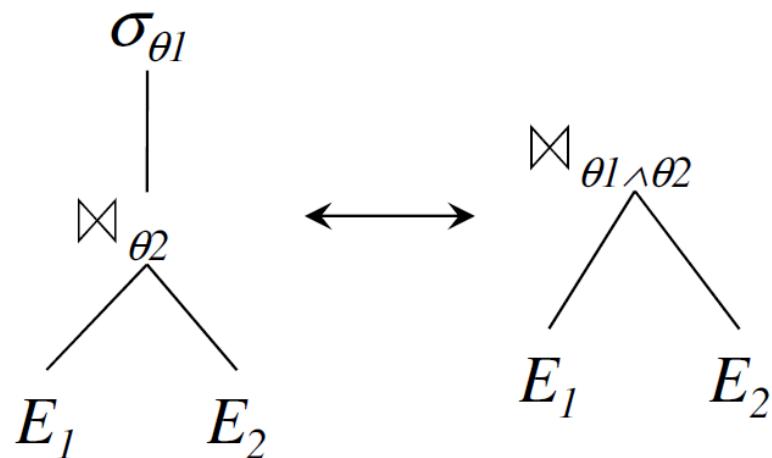
$$\Pi_{L_1}(\Pi_{L_2}(\dots(\Pi_{L_n}(E))\dots)) = \Pi_{L_1}(E)$$



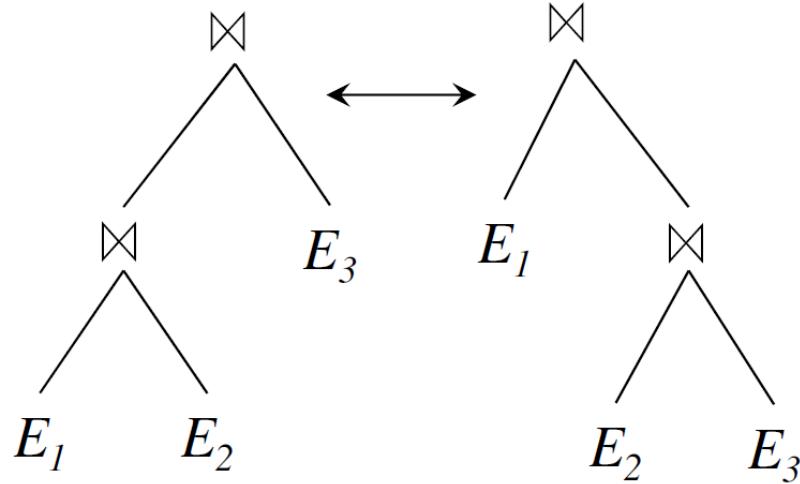
Joins are commutative

Selections can be combined
with joins

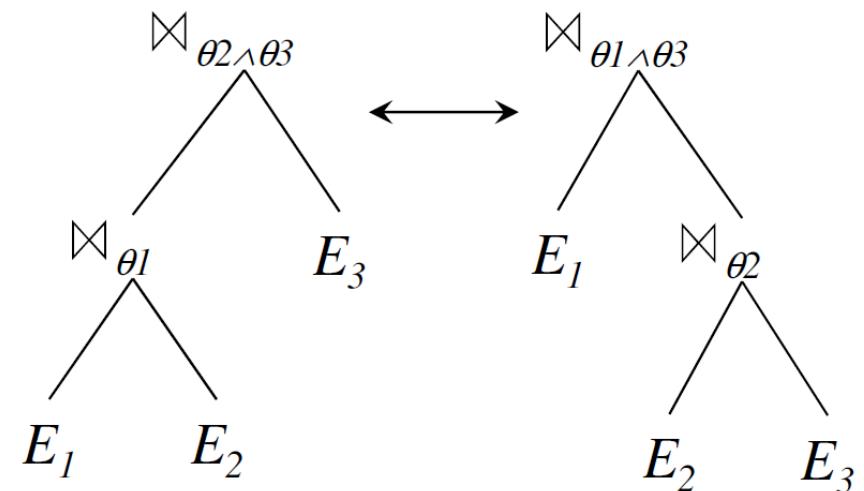
Joins are commutative



Joins are associative



when θ_2 only has attributes of E_2 and E_3



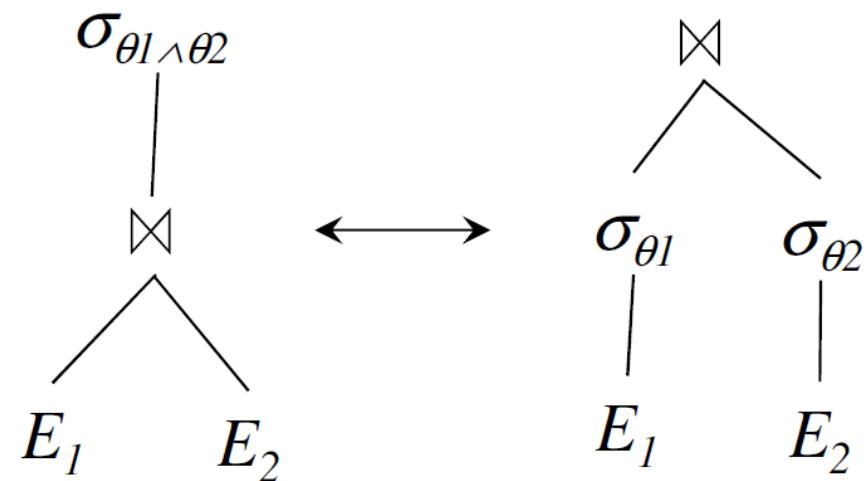
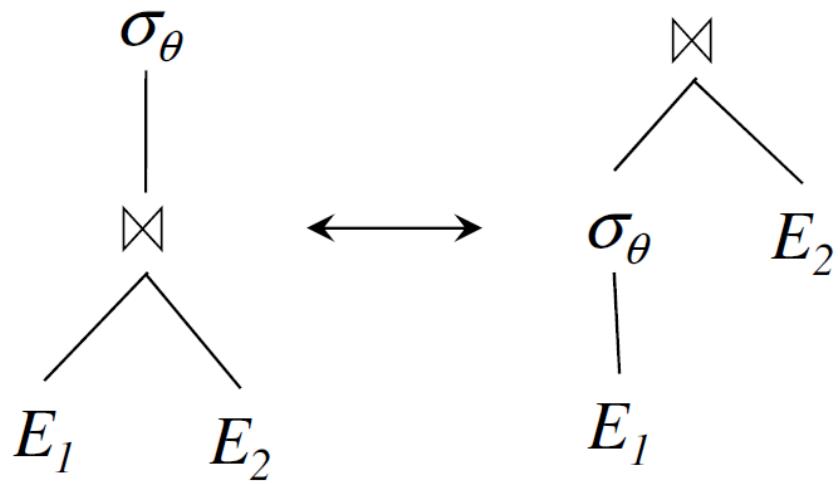


Selection Distributes over Joins



when θ only has attributes of E_1

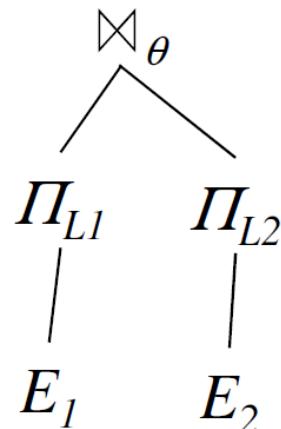
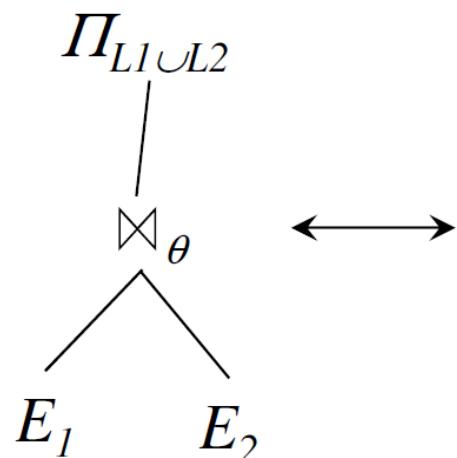
when θ_1 only has attributes of E_1 and
 θ_2 only has attributes of E_2



Projection Distributes over Join

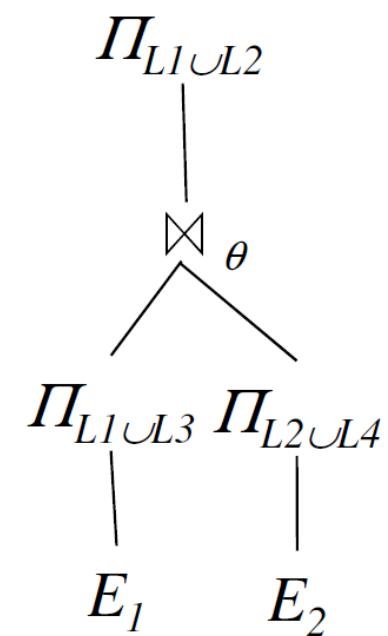
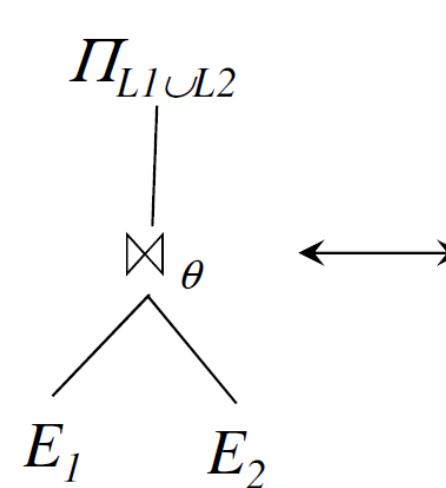


if θ only has attributes of $L_1 \cup L_2$



Let L_3 be attributes of E_1 that are in θ ,
but not in $L_1 \cup L_2$, and

Let L_4 be attributes of E_2 that are in θ ,
but not in $L_1 \cup L_2$

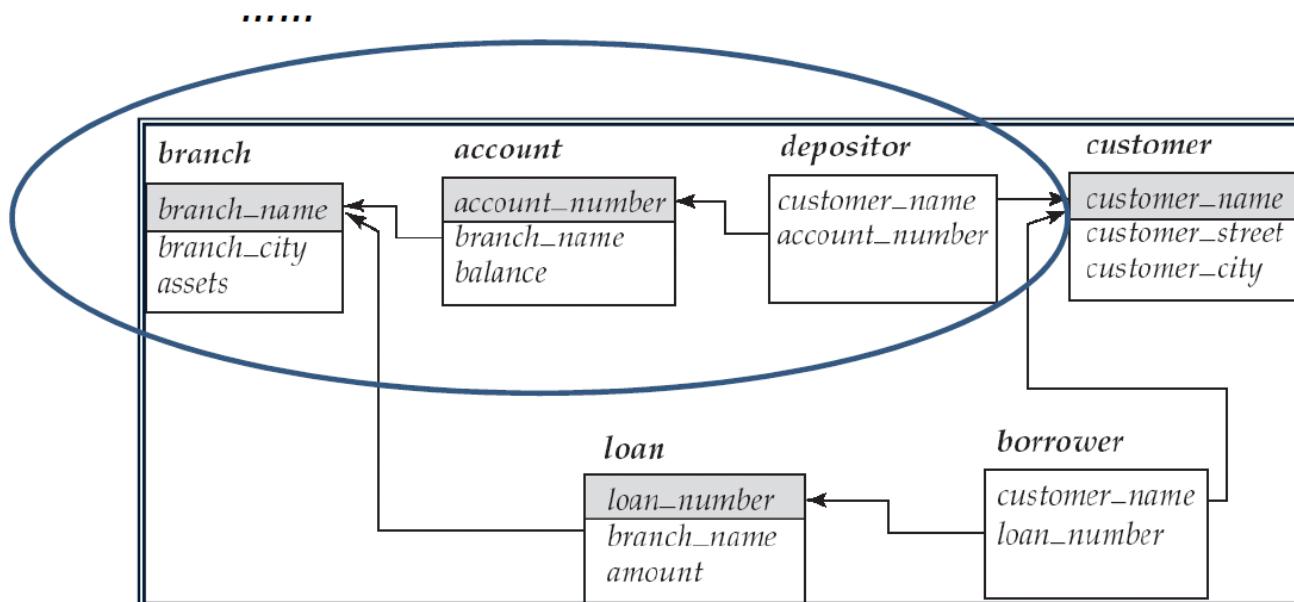


Example

branch (branch_name, branch_city, assets)

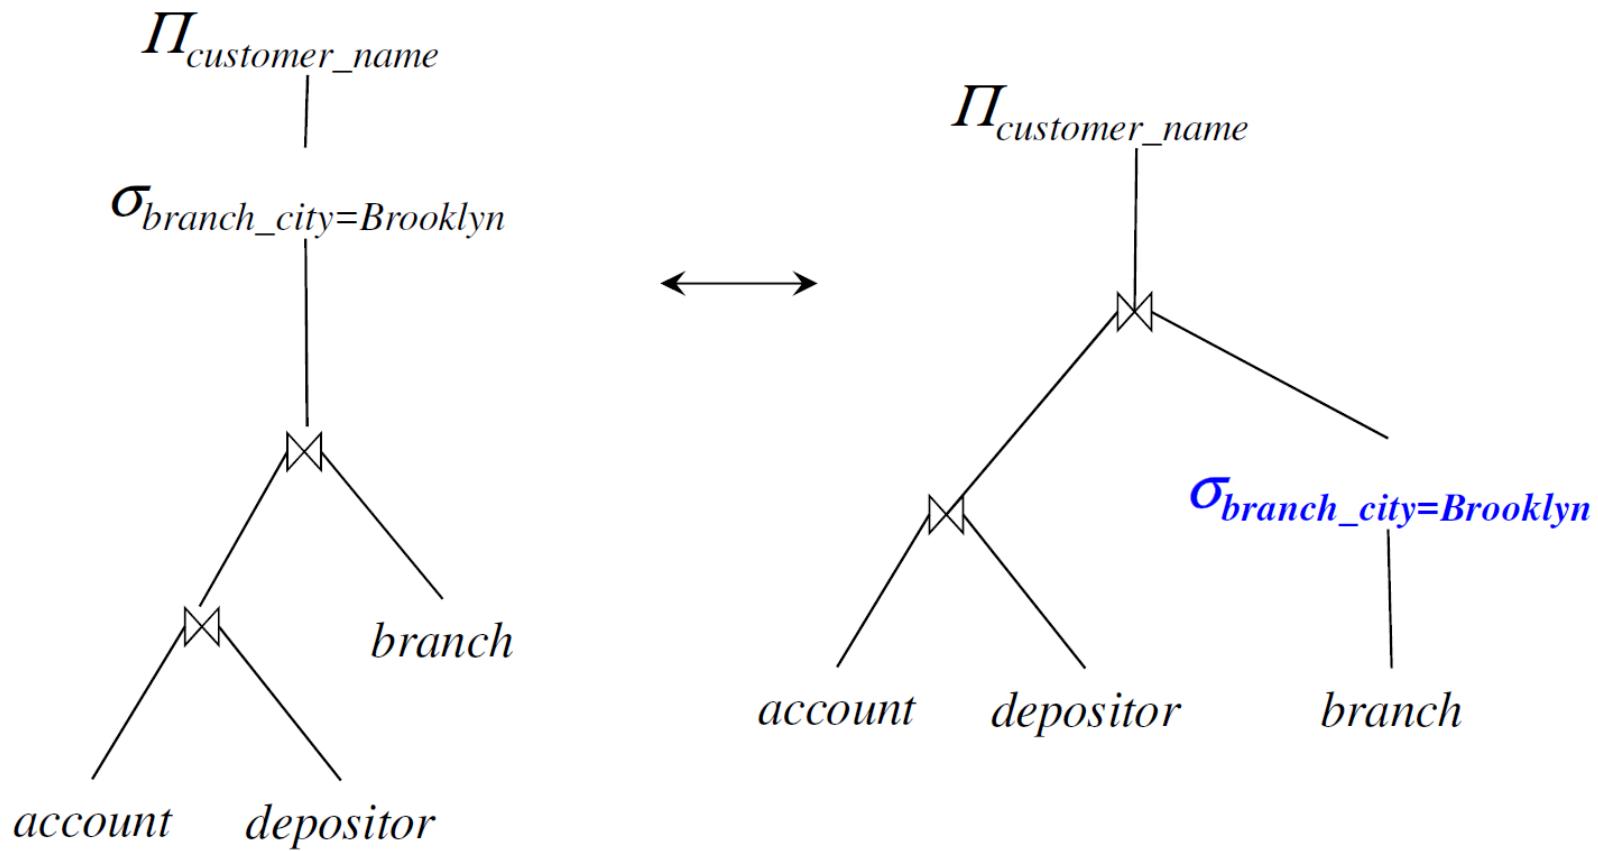
account (account_number, branch_name, balance)

depositor (customer_name, account_number)



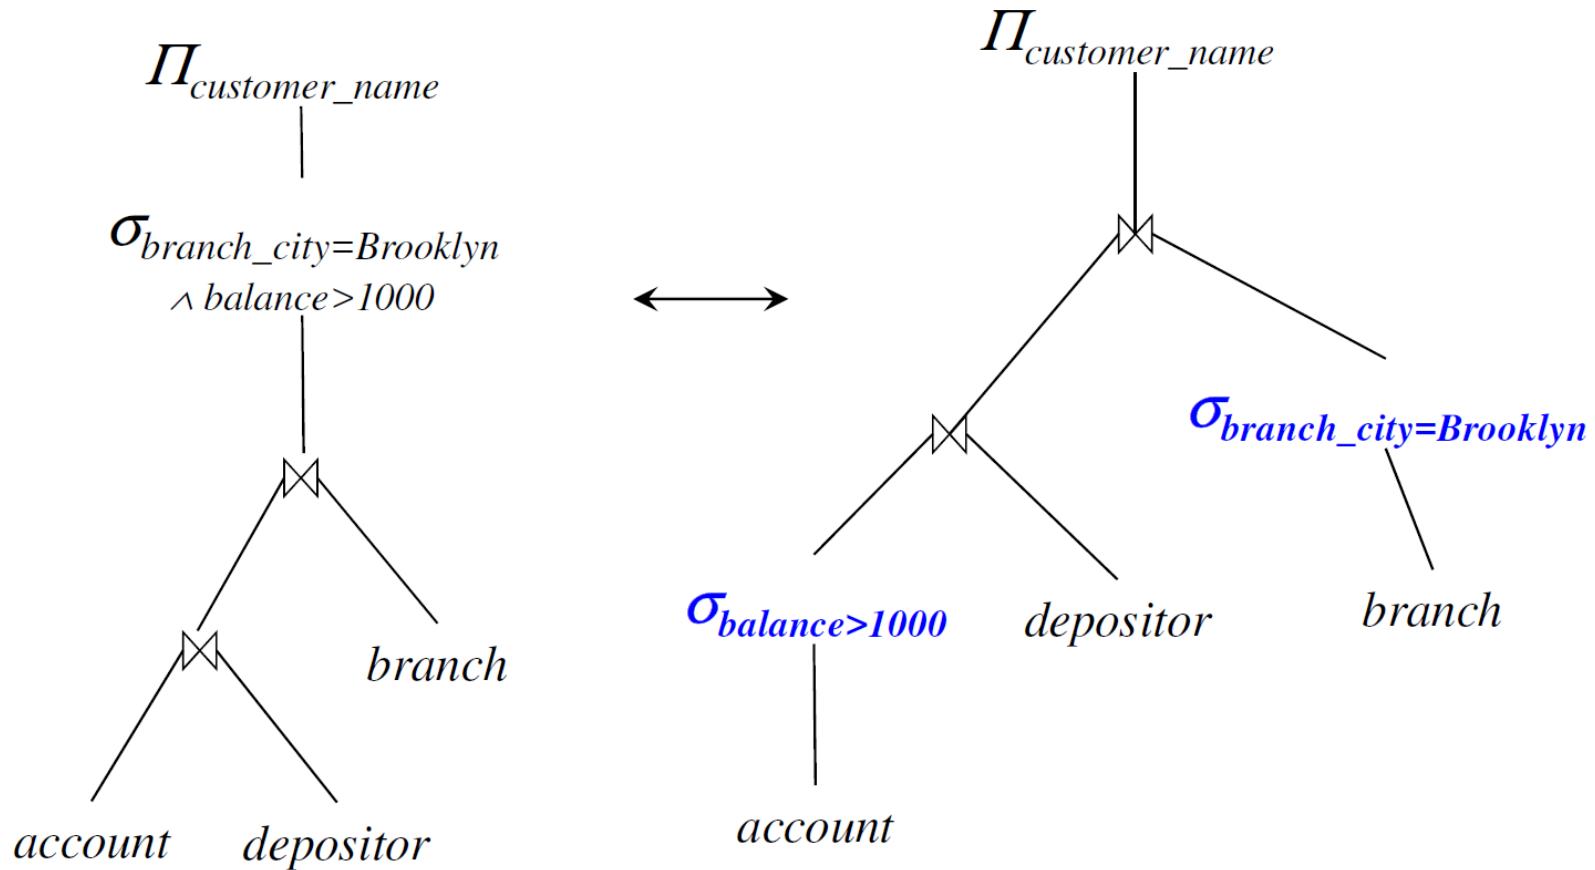
Reducing # of tuples early

- ◇ Query: Find the names of all customers who have an account at some branch located in Brooklyn



Reducing # of tuples early

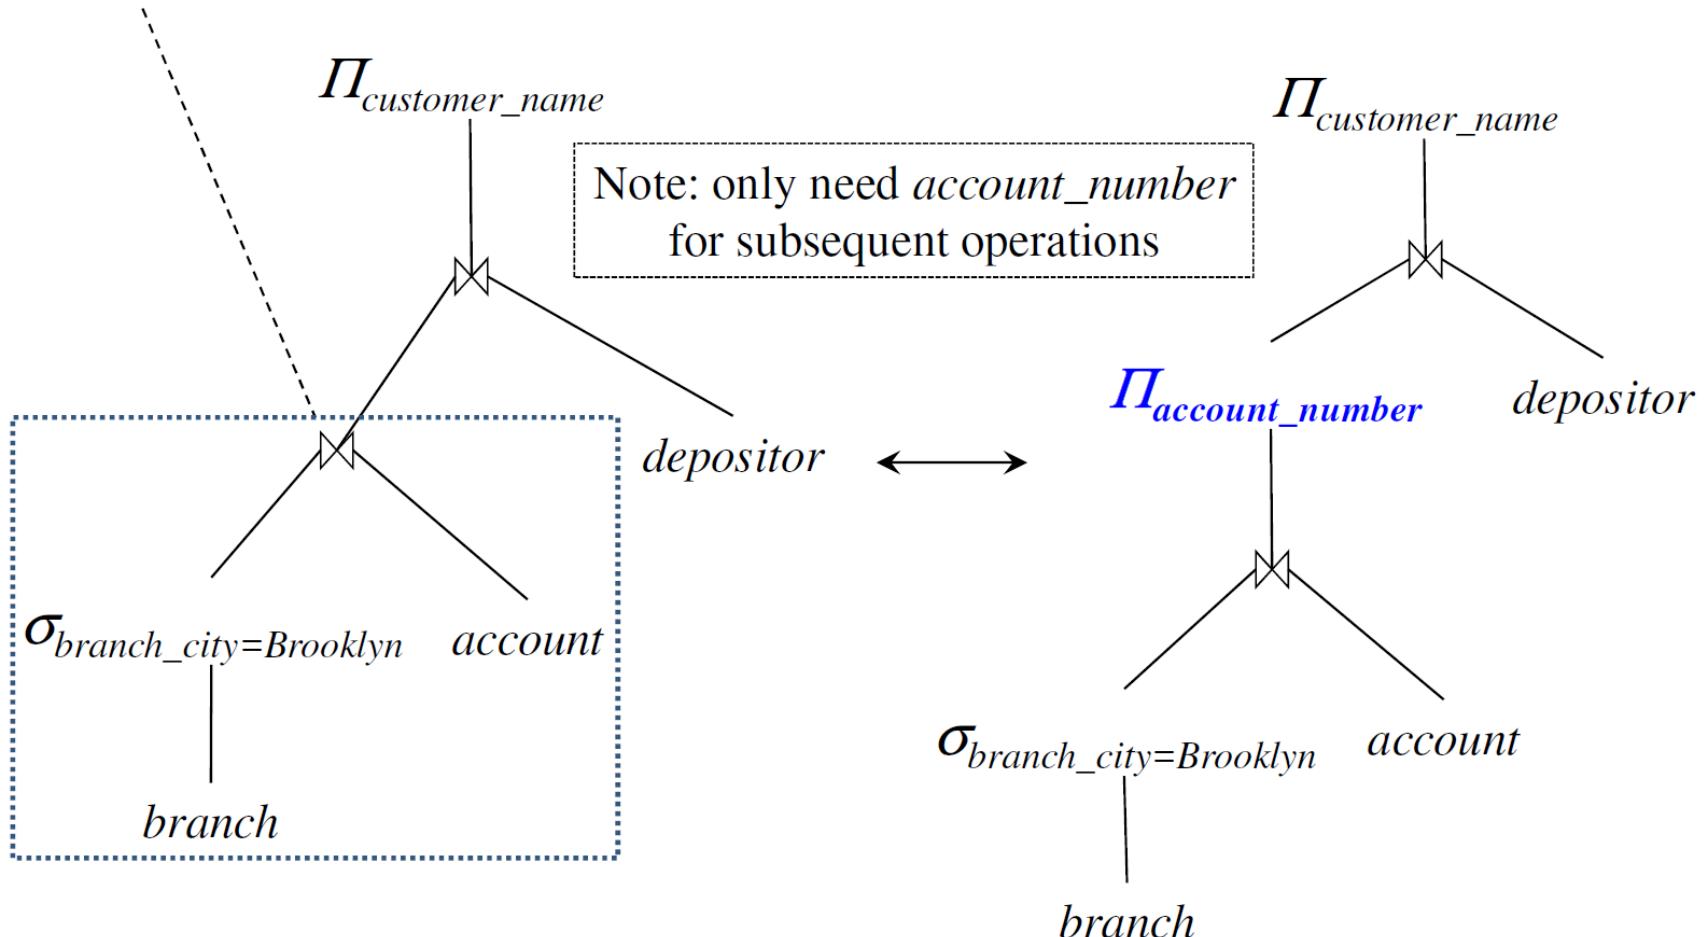
- ◆ Query: Find the names of all customers with an account at a Brooklyn branch whose account balance is over \$1000



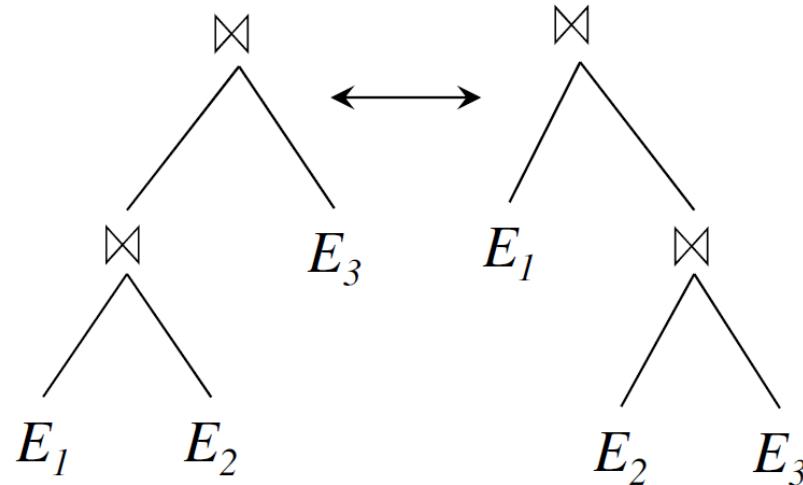
Reducing # of tuples early

the output relation schema here is:

$(branch_name, branch_city, assets, account_number, balance)$



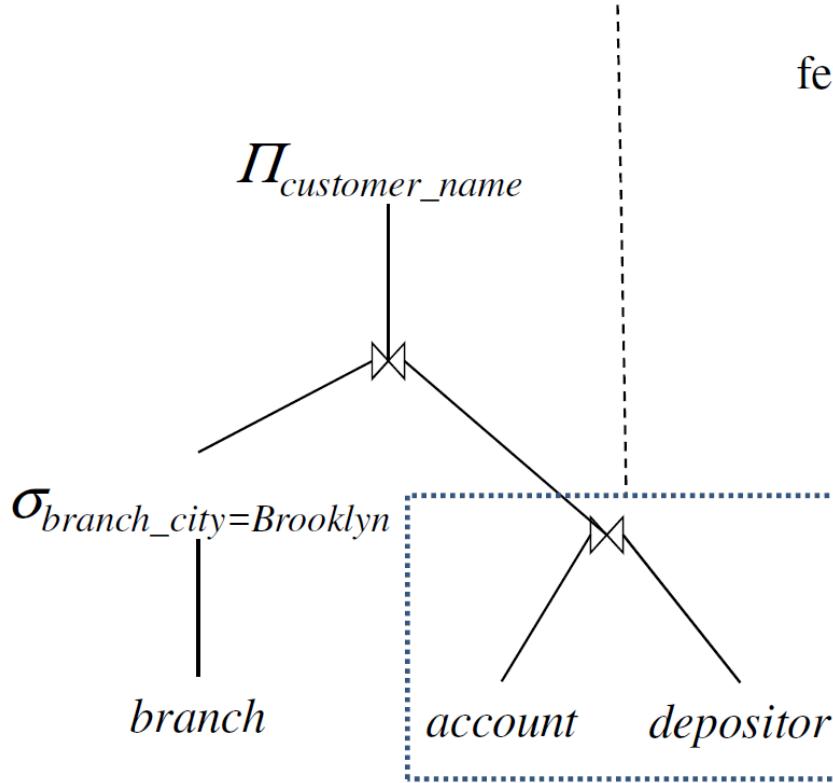
Join Ordering Example



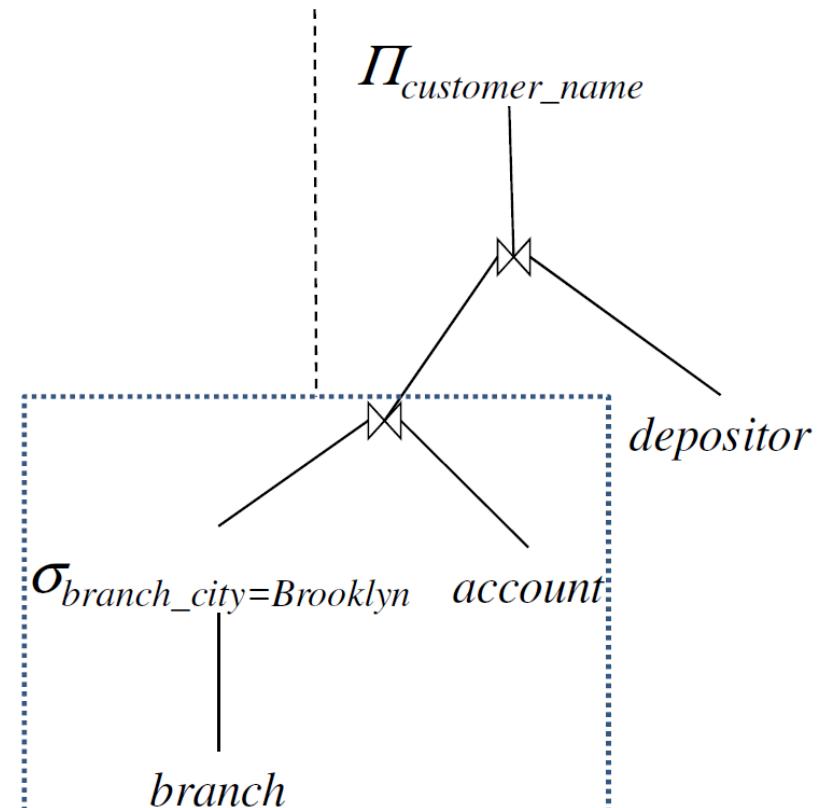
- ❖ If $E_2 \bowtie E_3$ is large and $E_1 \bowtie E_2$ is small, we choose
$$(E_1 \bowtie E_2) \bowtie E_3$$
so that we compute a smaller temporary relation

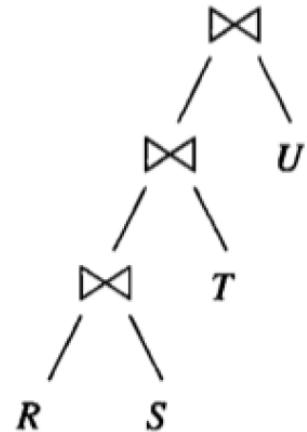
Join Ordering Example

this output relation is likely to be large



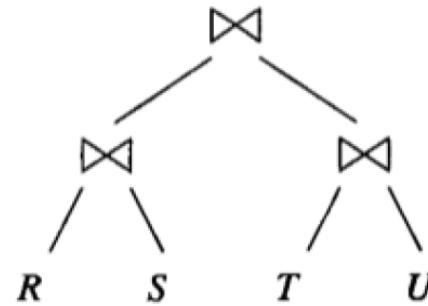
few accounts are in branches in Brooklyn
→ a smaller output relation





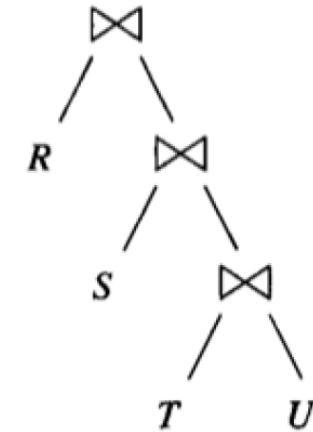
(a)

Left-deep join



(b)

Bushy



(c)

Right-deep



Dynamic Programming



- ❖ As a heuristic, optimizers enumerate left-deep plans only
 - ❖ Still $n!$ possible left-deep plans for n relations
- ❖ Dynamic programming approach
 - ❖ Find the best plans for 2 relations
 - ❖ Use them to find the best plans for 3 relations
 - ❖
 - ❖ Use them to find the best plans for n relations



Dynamic Programming Example

- ❖ Example from Chapter 16.6.4 of “Database Systems: The Complete Book”, 2nd Edition, Prentice Hall, 2009

- ❖ $R(a,b) \bowtie S(b,c) \bowtie T(c,d) \bowtie U(d,a)$
 - ❖ Each relation has 1000 tuples

- ❖ Cost estimation

- ❖ For simplicity, assume the cost of a join operator
= the # of tuples of input *intermediate relation(s)*
 - ◆ For accurate estimation, use the cost equations on **pages 10, 12**

- ❖ Estimate the output result size, between two relations r and s on join attribute x , by:
 - ◆ For accurate estimation,
use the equations on **page 45**

$$\frac{n_r * n_s}{\max\{V(r,x), V(s,x)\}}$$



Example



- ❖ 6 possible combinations: {R,S}, {R,T}, {R,U}, {S,T}, {S,U}, {T,U}
- ❖ Consider the combination {R,S}
 - ❖ Estimate its output size $R \bowtie S$
 $= 1000 * 1000 / \max(200, 100)$
 $= 5000$
 - ❖ Cost = 0 as input relations are not intermediate

$R(a, b)$	$S(b, c)$	$T(c, d)$	$U(d, a)$
$V(R, a) = 100$			$V(U, a) = 50$
$V(R, b) = 200$	$V(S, b) = 100$		
	$V(S, c) = 500$	$V(T, c) = 20$	
		$V(T, d) = 50$	$V(U, d) = 1000$

	{R, S}	{R, T}	{R, U}	{S, T}	{S, U}	{T, U}
Size	5000	1,000,000	10,000	2000	1,000,000	1000
Cost	0	0	0	0	0	0
Best plan	$R \bowtie S$	$R \bowtie T$	$R \bowtie U$	$S \bowtie T$	$S \bowtie U$	$T \bowtie U$

Example

$R(a, b)$	$S(b, c)$	$T(c, d)$	$U(d, a)$
$V(R, a) = 100$			$V(U, a) = 50$
$V(R, b) = 200$	$V(S, b) = 100$		
	$V(S, c) = 500$	$V(T, c) = 20$	
		$V(T, d) = 50$	$V(U, d) = 1000$

- ❖ 4 possible combinations: {R,S,T}, {R,S,U}, {R,T,U}, {S,T,U}
- ❖ Consider the combination {R,S,T}
 - ❖ Estimate its output size: $5000 * 1000 / \max(500, 20) = 10000$
 - ❖ Cost of $(R \bowtie S) \bowtie T = 5000$
 - ❖ Cost of $(S \bowtie T) \bowtie R = 2000$
 - ❖ Best plan: $(S \bowtie T) \bowtie R$

	{R, S}	{R, T}	{R, U}	{S, T}	{S, U}	{T, U}
Size	5000	1,000,000	10,000	2000	1,000,000	1000
Cost	0	0	0	0	0	0
Best plan	$R \bowtie S$	$R \bowtie T$	$R \bowtie U$	$S \bowtie T$	$S \bowtie U$	$T \bowtie U$

	{R, S, T}	{R, S, U}	{R, T, U}	{S, T, U}
Size	10,000	50,000	10,000	2,000
Cost	2,000	5,000	1,000	1,000
Best plan	$(S \bowtie T) \bowtie R$	$(R \bowtie S) \bowtie U$	$(T \bowtie U) \bowtie R$	$(T \bowtie U) \bowtie S$



Example



- ❖ Only 1 combination: $\{R, S, T, U\}$
 - ❖ The output size is the same for all cases below:
 - ❖ Cost of $((S \bowtie T) \bowtie R) \bowtie U = 2000 + 10000 = 12000$
 - ❖ Cost of $((R \bowtie S) \bowtie U) \bowtie T = 5000 + 50000 = 55000$
 - ❖ Cost of $((T \bowtie U) \bowtie R) \bowtie S = 1000 + 10000 = 11000$
 - ❖ Cost of $((T \bowtie U) \bowtie S) \bowtie R = 1000 + 2000 = 3000$
 - ❖ Best plan: $((T \bowtie U) \bowtie S) \bowtie R$

	$\{R, S, T\}$	$\{R, S, U\}$	$\{R, T, U\}$	$\{S, T, U\}$
Size	10,000	50,000	10,000	2,000
Cost	2,000	5,000	1,000	1,000
Best plan	$(S \bowtie T) \bowtie R$	$(R \bowtie S) \bowtie U$	$(T \bowtie U) \bowtie R$	$(T \bowtie U) \bowtie S$

Grouping	Cost
$((S \bowtie T) \bowtie R) \bowtie U$	12,000
$((R \bowtie S) \bowtie U) \bowtie T$	55,000
$((T \bowtie U) \bowtie R) \bowtie S$	11,000
$((T \bowtie U) \bowtie S) \bowtie R$	3,000

Example

- ❖ Cost of a plan = the sum of the cost of each operator
- ❖ Cost of each operator (see earlier slides)

- ❖ Need statistics of input relations

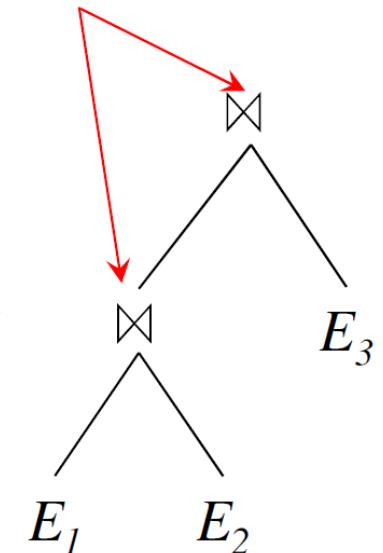
- ◆ E.g. number of tuples, sizes of tuples

- ❖ Some input can be another operator's output

- ◆ Need to estimate the size of such input

- ◆ To do so, we need additional statistics,

- e.g., number of distinct values for an attribute



- ❖ Notations

- ❖ n_r : number of tuples in a relation r

- ❖ f_r : blocking factor, i.e., number of tuples of r per block

- ❖ b_r : number of blocks in r

- ◆ $b_r = n_r / f_r$ if tuples are stored together physically in a file

- ❖ $V(r, A)$: number of distinct values for attribute A in r

Estimate Selection Size

- ◊ Let $size$ be the estimated number of tuples satisfying the condition
- ◊ Estimate the size of $\sigma_{A=v}(r)$
 - ◊ $size = 1$ if A is the primary key of r
 - ◊ $size = n_r / V(r, A)$: otherwise
- ◊ Estimate the size of $\sigma_{v_1 \leq A \leq v_2}(r)$
 - ◊ If $\min(r, A)$ and $\max(r, A)$ are available in the database catalog

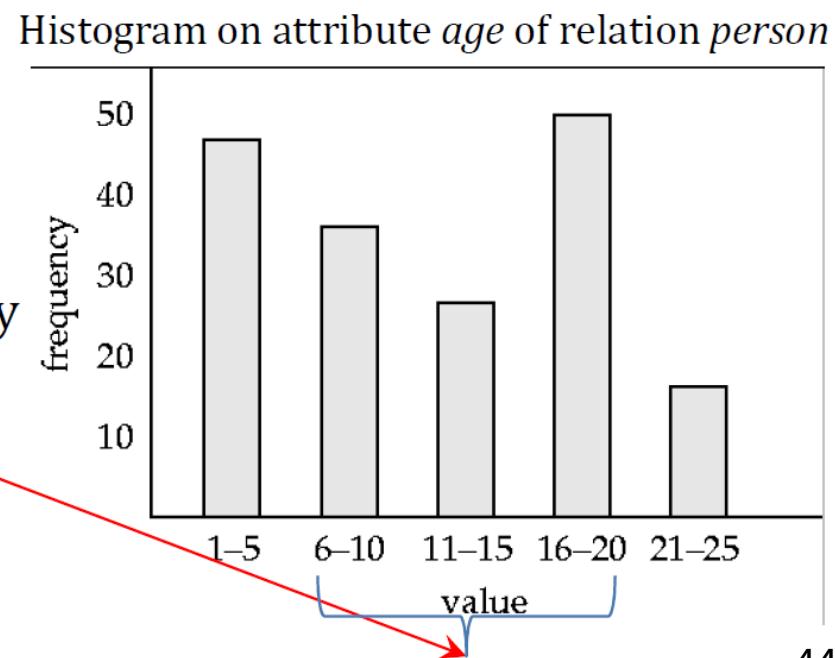
$$size = n_r \cdot \frac{v_2 - v_1}{\max(r, A) - \min(r, A)}$$

- ◊ If histograms available,
we can estimate the size more accurately
 - ◊ E.g., estimate $\sigma_{8 \leq \text{age} \leq 16}(\text{person})$

$$= 35 * (3/5) + 25 + 50 * (1/5)$$

$$= 21 + 25 + 10$$

$$= 56$$





Estimate Join Size



R : the schema of relation r

S : the schema of relation s

<i>Cases of attributes</i>	<i>Size of $r \bowtie s$</i>	<i>Remarks</i>
$R \cap S = \emptyset$	$n_r \times n_s$	
$R \cap S$ is a key for r	at most n_s	a tuple of s joins with at most one tuple from r
$R \cap S$ is a foreign key in s referencing r	exactly n_s	
$R \cap S = \{A\}$ is not a key for r nor s	$\frac{n_r * n_s}{\max\{V(r, A), V(s, A)\}}$	the lower estimate, assume that every tuple t in r (or s) produces tuples in $r \bowtie s$

Note: symmetric cases can be obtained by swapping two relations in the above table

- ◆ n_r : number of tuples in relation r
- ◆ $V(r, A)$: number of distinct values for attribute A in relation r



Example



Given information in the database catalog:

- ◆ $n_{customer} = 10000 \quad f_{customer} = 25$
→ $b_{customer} = 10000/25 = 400$
- ◆ $n_{depositor} = 5000 \quad f_{depositor} = 50$
→ $b_{depositor} = 5000/50 = 100$
- ◆ $V(depositor, customer_name) = 2500$
→ each customer has two accounts on average
- ◆ $customer_name$ in $depositor$ is a foreign key on $customer$
- ◆ $customer_name$ in $customer$ is a primary key
→ $V(customer, customer_name) = 10000$



Example



- ❖ Using foreign keys
 - ❖ *customer_name* in *depositor* is a foreign key on *customer*
 - ❖ Estimated result size
 - = $n_{depositor}$
 - = 5000 tuples
- ❖ Without using information about foreign keys:
 - ❖ $V(depositor, customer_name) = 2500$, and
 $V(customer, customer_name) = 10000$
 - ❖ Estimated size
 - = $5000 * 10000 / \max(2500, 1000)$
 - = $5000 * 10000 / 10000$
 - = 5000 tuples



Other factors



- ❖ Must consider the interaction of **physical operators** in a plan
 - ❖ Choosing the cheapest *physical operator* for each operation independently may not yield the best plan
- ❖ Examples
 - ❖ *Merge-join* may be more expensive than hash-join, but provides a *sorted output* which may reduce the cost for an *outer level aggregation*
 - ❖ *Nested-loop join* may enable *pipelining* (i.e., on-the-fly processing) which avoids writing intermediate relations on the disk



Summary



- ❖ Query processing
 - ❖ Query cost
 - ❖ Operations: selection, join
- ❖ Query optimization
 - ❖ Equivalence rules
 - ❖ Join ordering by dynamic programming
 - ❖ Cost and size estimation
- ❖ **Read this paper** (available in Blackboard):
 - ❖ “An overview of Query Optimization in Relational Systems”



谢谢！

DBGroup @ SUSTech

Dr. Bo Tang (唐博)

tangb3@sustech.edu.cn

