



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY




Distributed Databases I

南方科技大学
唐 博
tangb3@sustech.edu.cn

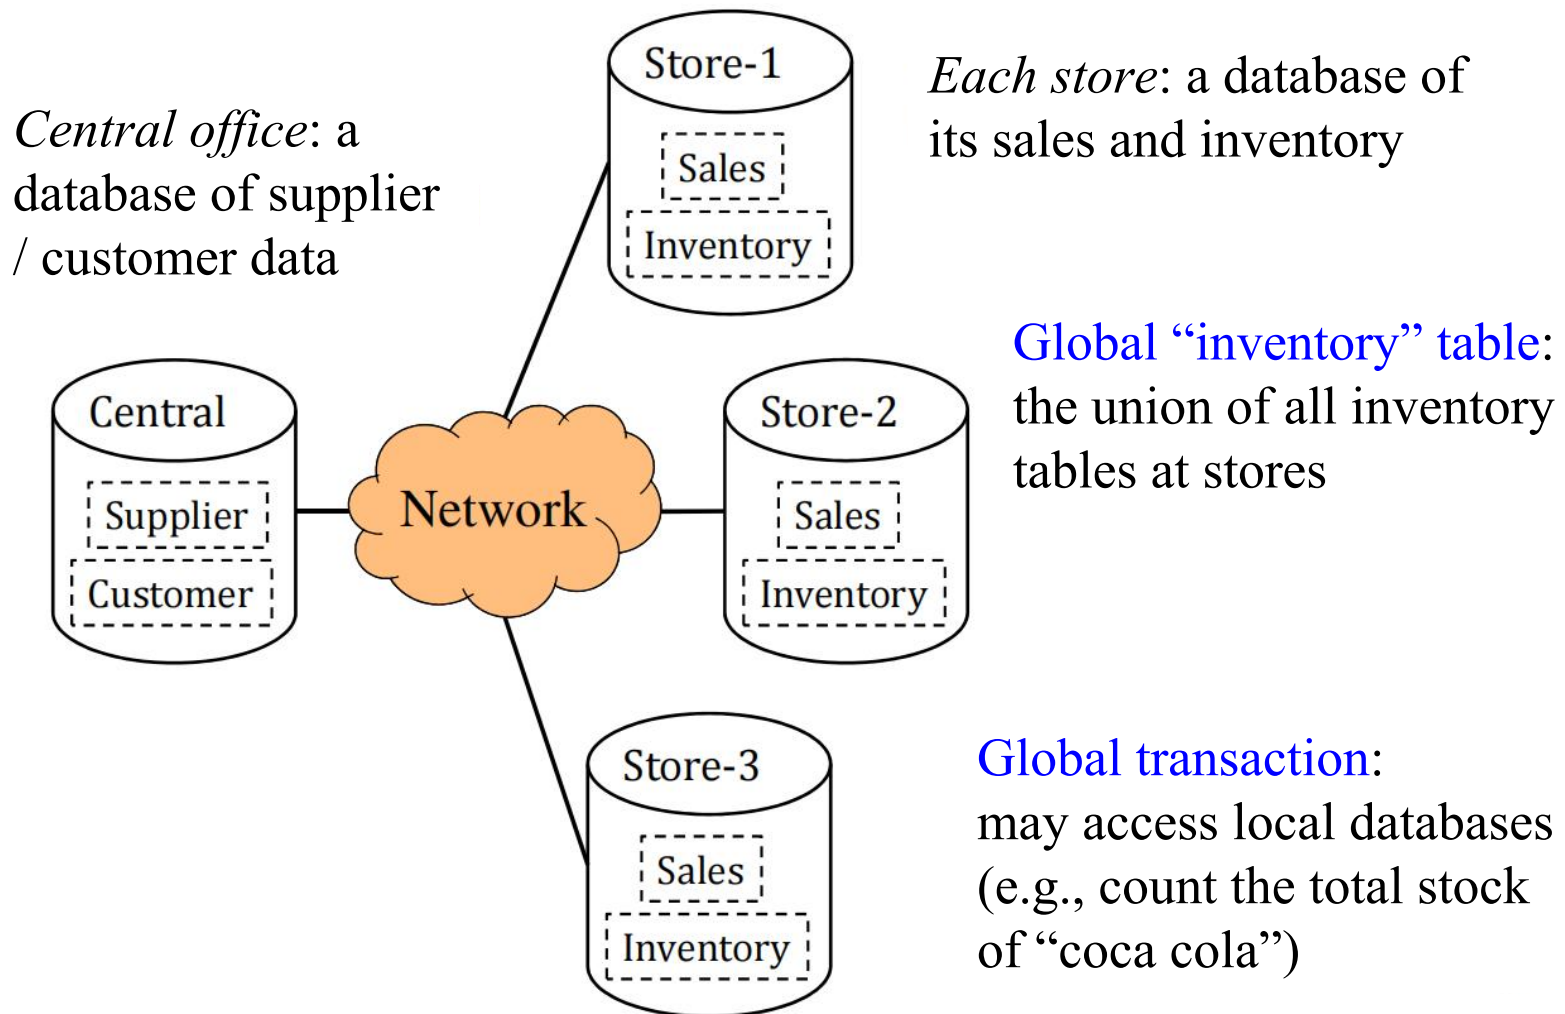




Lecture Objectives



- ❖ An overview of the distributed database
- ❖ Data replication and fragmentation
- ❖ The two-phase commit protocol





What is a distributed database?



- ❖ A collection of data with
 - ❖ *Distribution*: data are spread over different sites (of a network)
 - ❖ *Logical correlation*: data belong to the same system; some properties tie them together

Support *global transactions*:
accesses data at more than one site




Why distributed databases?



Management perspectives

- ❖ *Organizational requirement*: each division / branch (of the organization) may want to maintain its own DB
- ❖ *Interconnect existing DB's*: when multiple DBs already exist in an organization & need for global applications
- ❖ *Incremental growth*: support smooth incremental growth (e.g., adding a branch) with small impact on existing DBs






Why distributed databases?




Technical perspectives

- ❖ *Reduced communication overhead*: run sub-transactions at different sites ➡ sites transfer intermediate results (small) rather than entire tables
- ❖ *Parallel executions*: can execute some transactions in parallel at the participating sites
- ❖ *Reliability and availability*: can still run transactions despite failures of some sites




❖ Centralized DB



vs.

Distributed DB



❖ Revisit DBMS issues, how to:

❖ Store data?

[study today]

❖ By fragmentation, replication

❖ Ensure ACID properties?

❖ Recovery for A, D

[study today]

❖ By the two-phase commit protocol

❖ Concurrency for I

[next lecture]

❖ Process a query fast?


[next lecture]



Lecture Objectives



- ❖ An overview of the distributed database



- ❖ Data replication and fragmentation

- ❖ The two-phase commit protocol




Distributed Data Storage



- ❖ Assume relational data model
- ❖ *Fragmentation*
 - ❖ Relation is partitioned into several fragments stored at different sites
- ❖ *Replication*
 - ❖ System stores multiple copies of data at different sites
 - ❖ For faster retrieval and fault tolerance
- ❖ Replication and fragmentation can be combined
 - ❖ Relation is partitioned into several fragments
 - ❖ System stores several identical copies of each such fragment

- ❖ A relation (or fragment of a relation) is **replicated** if it is stored redundantly in two or more sites
- ❖ How to process **queries**?
 - ❖ Query either copy, OR
 - ❖ Query in parallel
- ❖ How to process **updates**?
 - ❖ Must update both copies?
 - ❖ We'll discuss more about this in the next lecture





Data Replication (Cont.)



❖ Advantages of Replication

- ❖ **Availability:** even when a site has failure, we can access copies at other sites
- ❖ **Parallelism:** queries on r may be processed by several nodes in parallel
- ❖ **Reduced data transfer:** relation r is available locally at each site containing a replica of r



Data Replication (Cont.)



❖ Disadvantages of Replication

- ❖ Expensive **updates**: each replica of relation r must be updated
- ❖ **Complex concurrency control**: updates to different replicas may cause inconsistent data
 - ❖ Need a concurrency control protocol for distributed DBs:
E.g., choose one copy as **primary copy** and apply concurrency control operations on primary copy

- ❖ Divide relation r into fragments r_1, r_2, \dots, r_n which contain *sufficient* information to reconstruct relation r

Example : relation account with schema

Account = (branch_name, customer_number, account_number, balance)

| <i>branch_name</i> | <i>customer_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|----------------------|-----------------------|----------------|
| Hillside | Lowman | A-305 | 500 |
| Hillside | Camp | A-226 | 336 |
| Valleyview | Camp | A-177 | 205 |
| Valleyview | Kahn | A-402 | 10000 |
| Hillside | Kahn | A-155 | 62 |
| Valleyview | Kahn | A-408 | 1123 |
| Valleyview | Green | A-639 | 750 |

Partition by Rows

Assign each tuple of r to one fragment


| <i>branch_name</i> | <i>customer_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|----------------------|-----------------------|----------------|
| Hillside | Lowman | A-305 | 500 |
| Hillside | Camp | A-226 | 336 |
| Hillside | Kahn | A-155 | 62 |

$$account_1 = \sigma_{branch_name = \text{"Hillside"}}(account)$$

| <i>branch_name</i> | <i>customer_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|----------------------|-----------------------|----------------|
| Valleyview | Camp | A-177 | 205 |
| Valleyview | Kahn | A-402 | 10000 |
| Valleyview | Kahn | A-408 | 1123 |
| Valleyview | Green | A-639 | 750 |


$$account_2 = \sigma_{branch_name = \text{"Valleyview"}}(account)$$

| <i>branch_name</i> | <i>customer_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|----------------------|-----------------------|----------------|
| Hillside | Lowman | A-305 | 500 |
| Hillside | Camp | A-226 | 336 |
| Hillside | Kahn | A-155 | 62 |



$$account_1 = \sigma_{branch_name = \text{"Hillside"}}(account)$$

| <i>branch_name</i> | <i>customer_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|----------------------|-----------------------|----------------|
| Valleyview | Camp | A-177 | 205 |
| Valleyview | Kahn | A-402 | 10000 |
| Valleyview | Kahn | A-408 | 1123 |
| Valleyview | Green | A-639 | 750 |



$$account_2 = \sigma_{branch_name = \text{"Valleyview"}}(account)$$

How to find out the sum of balance efficiently?

Partition by Columns

| <i>branch_name</i> | <i>customer_name</i> |
|--------------------|----------------------|
| Hillside | Lowman |
| Hillside | Camp |
| Valleyview | Camp |
| Valleyview | Kahn |
| Hillside | Kahn |
| Valleyview | Kahn |
| Valleyview | Green |

*Account =
(branch_name,
customer_number,
account_number,
balance)*

$$deposit_1 = \Pi_{branch_name, customer_name}(account)$$

| <i>account_number</i> | <i>balance</i> |
|-----------------------|----------------|
| A-305 | 500 |
| A-226 | 336 |
| A-177 | 205 |
| A-402 | 10000 |
| A-155 | 62 |
| A-408 | 1123 |
| A-639 | 750 |

Do we have **sufficient**
information to reconstruct
the original table?

$$deposit_2 = \Pi_{account_number, balance}(account)$$

- ❖ **Vertical fragmentation:** split the schema for relation r into several smaller schemas
- ❖ All schemas must contain a common candidate key (or superkey) to ensure lossless join property
- ❖ May need to add a special attribute (tuple-id) to each schema as a candidate key

$deposit_1 =$

$\Pi_{branch_name, customer_name, tuple_id}(account)$


| <i>branch_name</i> | <i>customer_name</i> | <i>tuple_id</i> |
|--------------------|----------------------|-----------------|
| Hillside | Lowman | 1 |
| Hillside | Camp | 2 |
| Valleyview | Camp | 3 |
| Valleyview | Kahn | 4 |
| Hillside | Kahn | 5 |
| Valleyview | Kahn | 6 |
| Valleyview | Green | 7 |

$deposit_2 =$


$\Pi_{account_number, balance, tuple_id}(account)$

| <i>account_number</i> | <i>balance</i> | <i>tuple_id</i> |
|-----------------------|----------------|-----------------|
| A-305 | 500 | 1 |
| A-226 | 336 | 2 |
| A-177 | 205 | 3 |
| A-402 | 10000 | 4 |
| A-155 | 62 | 5 |
| A-408 | 1123 | 6 |
| A-639 | 750 | 7 |

- ❖ Example: schema of an inventory relation
inventory(item-id, quantity, supplier, store)
- ❖ Horizontal Fragmentation



- ❖ Vertical Fragmentation





How to Improve Query Performance?

❖ Store together the **tuples** that are frequently accessed together

❖ E.g., likely to access tuples at the branch “Hillside” together

| <i>branch_name</i> | <i>customer_name</i> | <i>account_number</i> | <i>balance</i> |
|--------------------|----------------------|-----------------------|----------------|
| Hillside | Lowman | A-305 | 500 |
| Hillside | Camp | A-226 | 336 |
| Hillside | Kahn | A-155 | 62 |

❖ Store together the **attributes** that are frequently accessed together

❖ E.g., likely to access the attributes account number and balance together

| <i>account_number</i> | <i>balance</i> | <i>tuple_id</i> |
|-----------------------|----------------|-----------------|
| A-305 | 500 | 1 |
| A-226 | 336 | 2 |
| A-177 | 205 | 3 |
| | | |



How to Improve Query Performance?



- ❖ Different transactions may access data with different access patterns
- ❖ Difficult to decide the fragmentation manually (by DB administrator)
- ❖ Any **automatic** method for this problem?
 - ❖ First, extract access patterns from transactions
 - ❖ Then, design the fragmentation accordingly

Use attribute usage to derive a good vertical fragmentation

| Attribute usage matrix | | | | | | | | | | Type | Number of accesses per time period | |
|----------------------------|---|---|---|---|---|---|---|---|---|------|------------------------------------|------------|
| Attributes Transactions | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | | |
| T1 | 1 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | R | Acc 1 = 25 |
| T2 | 0 | 1 | 1 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | R | Acc 2 = 50 |
| T3 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 0 | 1 | R | Acc 3 = 25 |
| T4 | 0 | 1 | 0 | 0 | 0 | 0 | 1 | 1 | 0 | 0 | R | Acc 4 = 35 |
| T5 | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 1 | 1 | 0 | U | Acc 5 = 25 |
| T6 | 1 | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | U | Acc 6 = 25 |
| T7 | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | U | Acc 7 = 25 |
| T8 | 0 | 0 | 1 | 1 | 0 | 1 | 0 | 0 | 1 | 1 | U | Acc 8 = 15 |

Fig.1 Attribute usage matrix

| Attributes | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
|------------|----|-----|-----|----|----|----|----|-----|-----|----|
| 1 | 75 | 25 | 25 | 0 | 75 | 0 | 50 | 25 | 25 | 0 |
| 2 | 25 | 110 | 75 | 0 | 25 | 0 | 60 | 110 | 75 | 0 |
| 3 | 25 | 75 | 115 | 15 | 25 | 15 | 25 | 75 | 115 | 15 |
| 4 | 0 | 0 | 15 | 40 | 0 | 40 | 0 | 0 | 15 | 40 |
| 5 | 75 | 25 | 25 | 0 | 75 | 0 | 50 | 25 | 25 | 0 |
| 6 | 0 | 0 | 15 | 40 | 0 | 40 | 0 | 0 | 15 | 40 |
| 7 | 50 | 60 | 25 | 0 | 50 | 0 | 85 | 60 | 25 | 0 |
| 8 | 25 | 110 | 75 | 0 | 25 | 0 | 60 | 110 | 75 | 0 |
| 9 | 25 | 75 | 115 | 15 | 25 | 15 | 25 | 75 | 115 | 15 |
| 10 | 0 | 0 | 15 | 40 | 0 | 40 | 0 | 0 | 15 | 40 |

Fig.2 Attribute affinity (AA) matrix

Example adapted from the paper:
“Vertical Partitioning for Database Design: A Graphical Algorithm”. SIGMOD 1989.



❖ Horizontal:

- ❖ Allows parallel processing (on fragments with different tuples)

❖ Vertical:

- ❖ Allows parallel processing (on fragments with different attributes)
- ❖ Tuple-id attribute allows efficient joining of vertical fragments
- ❖ Vertical and horizontal fragmentation can be mixed
 - ❖ Fragments may be further fragmented to an arbitrary depth

- ❖ An overview of the distributed database

- ❖ Data replication and fragmentation



- ❖ The two-phase commit protocol

This protocol aims to achieve the properties 'A' and 'D'
Don't confuse it with the 2PL protocol (for property 'I')

❖ Hard to achieve all three properties together

❖ C: Consistency

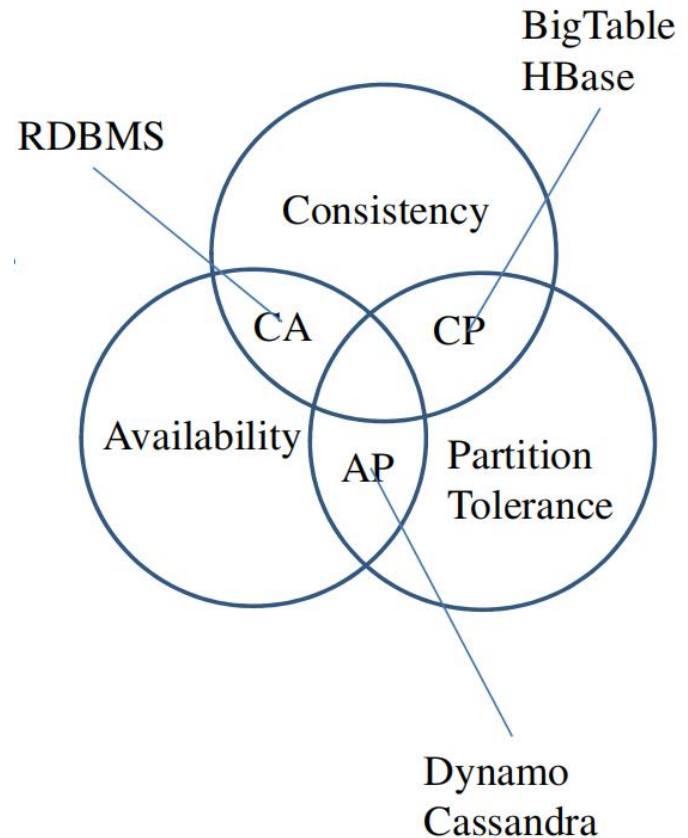
❖ All users can access the up-to-date copy of the data

❖ A: Availability


❖ The system can work properly even with node failures

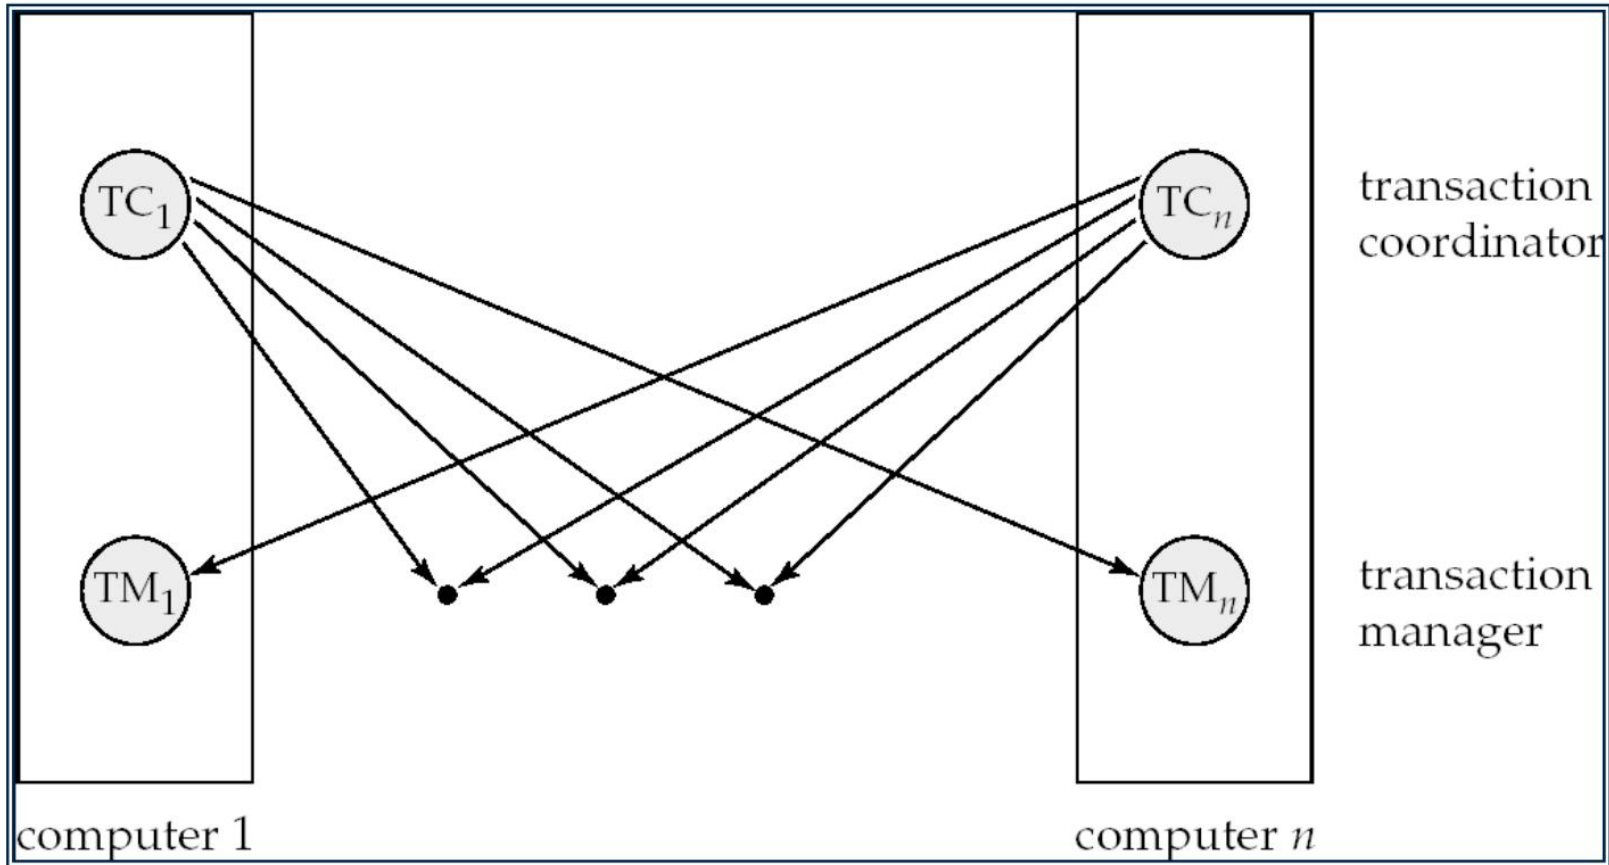
❖ P: Partitioning tolerance

❖ The system can work properly even with network/message failures



- ❖ Example: consider replicated data at two sites
 - ❖ Consistency: When we update data in S1, need to replicate this update in S2
 - ❖ Availability: When a site is running, we can query/update data from it (via network)
 - ❖ Partition tolerance: If the network fails, we can still query/update the local site
- ❖ When the network fails
 - ❖ Allow both sites available ➡ data may not be up-to-date
 - ❖ Keep consistency ➡ cannot make both sites available





- ❖ Transaction may access data at several sites
- ❖ Each site has a local transaction manager to:
 - ❖ Maintain a **log** for recovery purposes
 - ❖ Participate in the concurrent execution of transactions at that site
- ❖ Each site has a transaction coordinator to:
 - ❖ Start the execution of **global transactions** that originate at the site
 - ❖ Distribute **sub-transactions** to appropriate sites
 - ❖ Coordinate the termination of each transaction that originates at the site, which may result in:
commit the transaction at all sites / abort at all sites



System Failure Modes



- ❖ Failures unique to distributed systems:
 - ❖ Message loss
 - ❖ Handled by network protocols (e.g., TCP-IP)
- ❖ Communication link failure
 - ❖ Handled by network protocols, by routing messages via alternative links
- ❖ Site failure
- ❖ **Network partition**
 - ❖ It happens when the network is split into subsystems that lack connection
 - ❖ Note: a subsystem may consist of a single node


- ❖ Commit protocols ensure **atomicity** across sites
 - ❖ a transaction (which executes at multiple sites) must either be committed at all sites, OR aborted at all sites.

not acceptable to have a transaction
committed at one site and aborted at another

- ❖ The *two-phase commit* (2PC) protocol is widely used
 - ❖ It ensures atomicity property despite network / site failures
 - ❖ Suppose that each site uses recovery protocol to ensure sub-transaction atomicity



Two-phase commit: Notations



❖ T : a global transaction

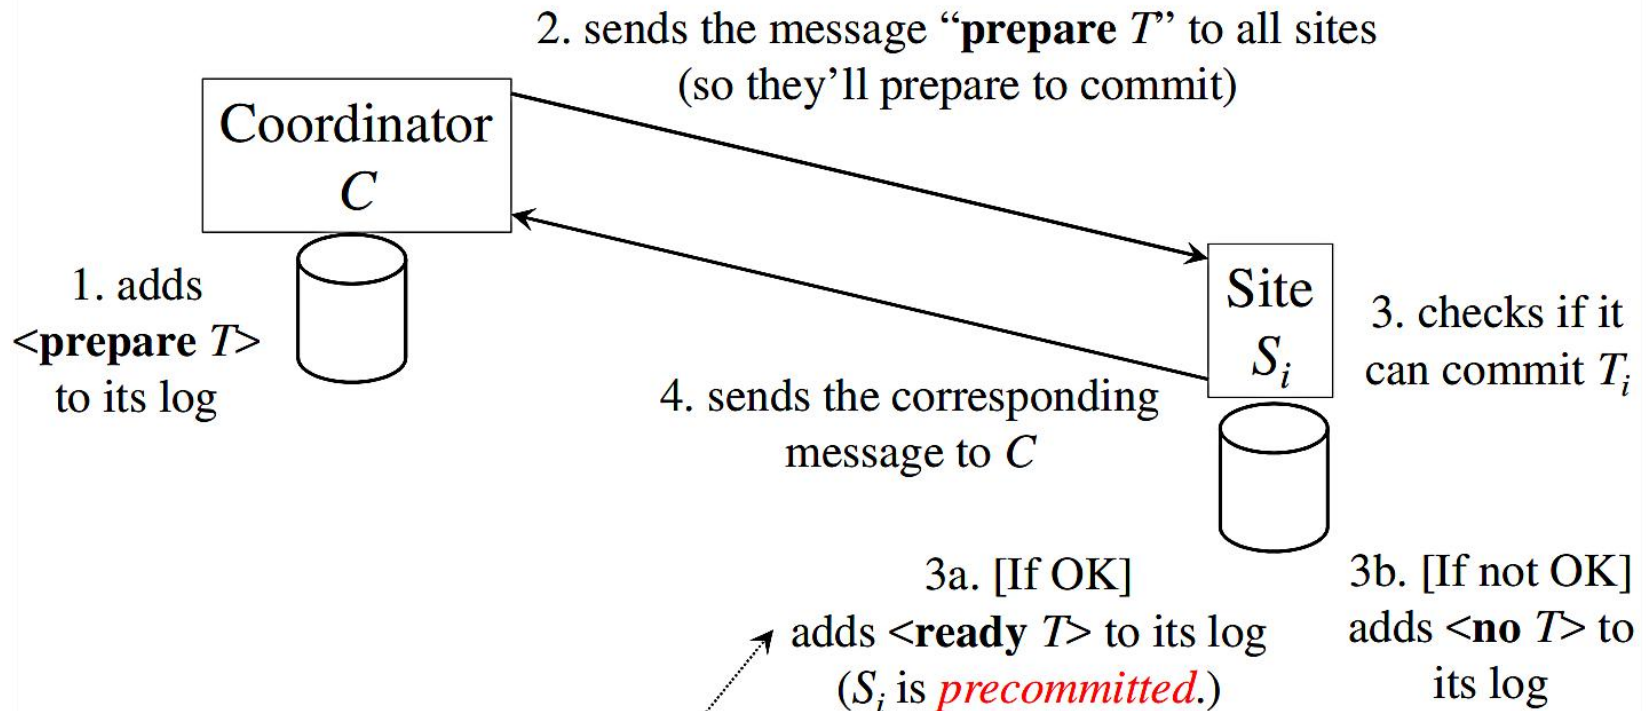
❖ T_1, T_2, \dots, T_n : sub-transactions of T

❖ T_i will be executed at participant sites S_i

❖ C : the coordinator of T


❖ This site monitors the sub-transactions and decides whether T should commit or not

(when all sub-transactions finish)



Q: Why S_i needs to write $\langle \text{ready } T \rangle$ to its stable log?


A: If failures occur later (before the protocol finishes), then it can commit/abort T_i



Cases for the coordinator C

Actions for C

| | |
|---|--|
| (1): received a “ ready T ” from all sites | <p><u>decides to commit T</u></p> <ul style="list-style-type: none"> • adds <commit T> to its log • sends the message “commit T” to all sites |
| (2): received an “ abort T ” from some sites <hr/> (3): not heard from some site S_i after a certain timeout period [C assumes S_i is down] | <p><u>decides to abort T</u></p> <ul style="list-style-type: none"> • adds <abort T> to its log • sends the message “abort T” to all sites |

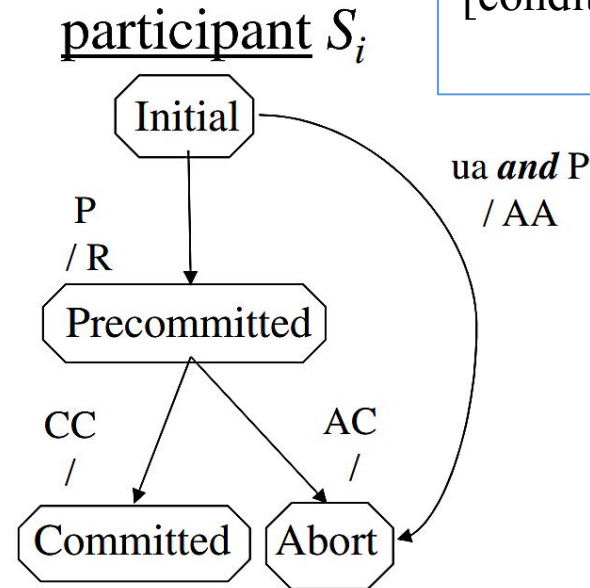



For each participant S_i

- ❖ Received “**commit** T ”:
 - ❖ it adds **<commit T >** to its log and commits T_i locally
- ❖ Received “**abort** T ”:
 - ❖ it adds **<abort T >** to its log and aborts T_i locally

Format

[condition, message received]
/ [message sent]



◆ Conditions:

- ◆ ua = unilateral abort
- ◆ tm = timeout

◆ Messages:

- ◆ P = "prepare T"
- ◆ R = "ready T"
- ◆ AA = "abort T" answer
- ◆ AC = "abort T" command
- ◆ CC = "commit T" command



Recovery




- ❖ No failure ➡ 2PC ensures that either all sub-transactions **commit** or all of them **abort**
- ❖ Site / network failure ➡ need to make sure that the site's recovery is consistent with the global decision for T
- ❖ Types of failure:
 - ❖ (1) a site failure
 - ❖ (2) a coordinator failure
 - ❖ (3) network partition failure

(1) Site Failure

- ❖ When a site S_i recovers after a failure, it checks its log for entries for T . If it finds:
 - ❖ **<commit T >**: T has committed, **redo** T_i
 - ❖ **<abort T >**: T has been aborted, **undo** T_i
 - ❖ **<no T >**: S_i has not received the decision from C yet, but the decision must be to abort T , so **undo** T_i
 - ❖ **<ready T >**: S_i does not know the decision. It asks C to determine whether T has committed or aborted.
 - ❖ None of the above: C could not have decided to commit T . It would be safe to abort T_i (so **undo** T_i)


(2) Coordinator Failure



- ❖ If the coordinator fails and all the “living” participants are precommitted ($\langle \text{ready } T \rangle$ logged), no one knows the final decision until the coordinator recovers.
 - ❖ All “living” participants are blocked
 - ❖ This is the *blocking problem* of 2PC
- ❖ The blocking problem of 2PC is undesirable
 - ❖ sub-transactions are holding locks on data items
 - ❖ cause severe blocking to other transactions in the system

(3) Network Partition Failure

- ❖ When a network partition occurs, a participant cannot communicate with the coordinator
 - ❖ In that case, the separated participant assumes the coordinator fails, and
 - ❖ The coordinator assumes that the separated participant fails
- ❖ So all sites execute the same 2PC protocol





Summary



- ❖ An overview of the distributed database
- ❖ Data replication and fragmentation
- ❖ The two-phase commit protocol

Readings after the class

- ❖ Chapters 3 and 12.4 in the book
Ozsu, and Valuriez. Principles of Distributed Database System,
3rd Ed, Springer, 2011 (free online)



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



谢谢!

DBGGroup @ SUSTech
Dr. Bo Tang (唐博)
tangb3@sustech.edu.cn

